


Experiment - 1

Bank Churn ANN


```
import numpy as np
import pandas as pd
```

```
churn_data = pd.read_csv('/workspace/FDP/Day3/Churn_Modelling.csv', delimiter = ',')
churn_data.head(5)
```




	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1

```
churn_data.columns
```



```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
churn_data = churn_data.set_index('RowNumber')
churn_data.head()
```




	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
RowNumber												
1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	
2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	
3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	
4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	
5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	

```
churn_data.shape
```



```
(10000, 13)
```

```
churn_data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 1 to 10000
Data columns (total 13 columns):
#   Column             Non-Null Count  Dtype
---  -
0   CustomerId         10000 non-null  int64
1   Surname            10000 non-null  object
2   CreditScore        10000 non-null  int64
3   Geography          10000 non-null  object
4   Gender             10000 non-null  object
5   Age               10000 non-null  int64
6   Tenure            10000 non-null  int64
7   Balance           10000 non-null  float64
8   NumOfProducts     10000 non-null  int64
9   HasCrCard         10000 non-null  int64
10  IsActiveMember    10000 non-null  int64
11  EstimatedSalary    10000 non-null  float64
12  Exited            10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1.1+ MB
```

```
churn_data.info()
```

```
churn_data.isna().sum()
```

```
CustomerId      0
Surname          0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
churn_data.nunique()
```

```
CustomerId      10000
Surname         2932
CreditScore     460
Geography       3
Gender          2
Age             70
Tenure         11
Balance        6382
NumOfProducts   4
HasCrCard       2
IsActiveMember  2
EstimatedSalary 9999
Exited          2
dtype: int64
```

```
churn_data.drop(['CustomerId', 'Surname'],axis=1,inplace=True)
```

```
churn_data.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber											
1	619	France	Female	42	2	0.00	1	1	1	101348.88	1
2	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
3	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
4	699	France	Female	39	1	0.00	2	0	0	93826.63	0
5	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
churn_data.shape
```

```
(10000, 11)
```

```
from matplotlib import pyplot as plt
import seaborn as sns
from scipy import stats
df = churn_data.copy()
```

```
def plot_univariate(col):
    if(df[col].nunique()>2):
        plt.figure(figsize=(10,7))
        h = 0.15
        rot=90
    else:
        plt.figure(figsize=(6,6))
        h = 0.5
        rot=0
    plot = sns.countplot(x = df[col], palette='pastel')

    for bars in plot.containers:
        for p in bars:
            plot.annotate(format(p.get_height()), (p.get_x() + p.get_width()*0.5, p.get_height()),
                           ha = 'center', va = 'bottom')
            plot.annotate(f'{p.get_height()*100/df[col].shape[0] : .1f}%', (p.get_x() + p.get_width()*0.5, h*p.get_height()),
                           ha = 'center', va = 'bottom', rotation=rot)
```

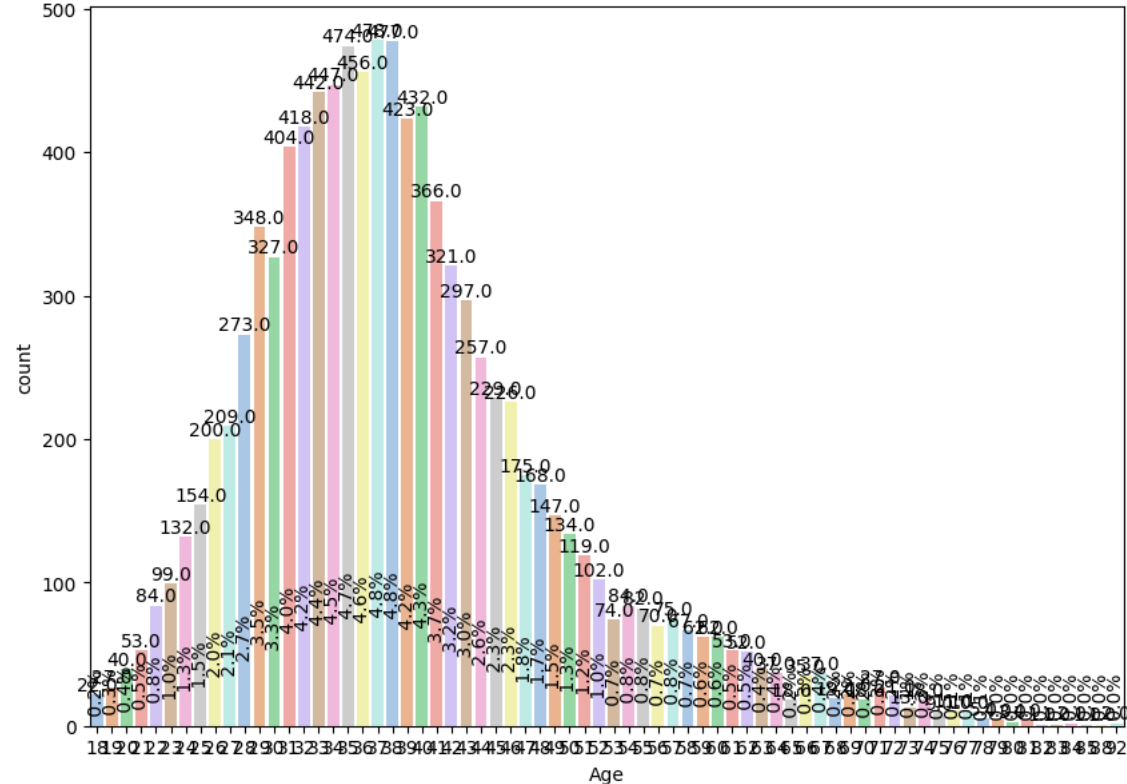
```
def spearman(df,hue):
    feature = []
    correlation = []
```

```
result = []
for col in df.columns:
    corr, p = stats.spearmanr(df[col], df[hue])
    feature.append(col)
    correlation.append(corr)
    alpha = 0.05
    if p > alpha:
        result.append('No correlation (fail to reject H0)')
    else:
        result.append('Some correlation (reject H0)')
c = pd.DataFrame({'Feature Name':feature,'correlation coefficient':correlation, 'Inference':result})
display(c)
```

```
plot_univariate('Age')
```

```

/tmp/ipykernel_6944/4200767248.py:10: FutureWarning:
    Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le
    plot = sns.countplot(x = df[col], palette='pastel')
```



```
spearman(churn_data, 'Age')
```

	Feature Name	correlation coefficient	Inference
0	CreditScore	-0.007974	No correlation (fail to reject H0)
1	Geography	0.035351	Some correlation (reject H0)
2	Gender	-0.029785	Some correlation (reject H0)
3	Age	1.000000	Some correlation (reject H0)
4	Tenure	-0.010405	No correlation (fail to reject H0)
5	Balance	0.033304	Some correlation (reject H0)
6	NumOfProducts	-0.058566	Some correlation (reject H0)
7	HasCrCard	-0.015278	No correlation (fail to reject H0)
8	IsActiveMember	0.039839	Some correlation (reject H0)
9	EstimatedSalary	-0.002431	No correlation (fail to reject H0)
10	Exited	0.323968	Some correlation (reject H0)

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
churn_data[['Geography', 'Gender']] = churn_data[['Geography', 'Gender']].apply(le.fit_transform)
```

```
churn_data.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber											
1	619	0	0	42	2	0.00	1	1	1	101348.88	1
2	608	2	0	41	1	83807.86	1	0	1	112542.58	0
3	502	0	0	42	8	159660.80	3	1	0	113931.57	1
4	699	0	0	39	1	0.00	2	0	0	93826.63	0
5	850	2	0	43	2	125510.82	1	1	1	79084.10	0

```
y = churn_data.Exited
X = churn_data.drop(['Exited'],axis=1)
```

```
X.columns
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary'],
      dtype='object')
```

```
y
```

```
RowNumber
1         1
2         0
3         1
4         0
5         0
..
9996      0
9997      0
9998      1
9999      1
10000     0
Name: Exited, Length: 10000, dtype: int64
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 2)
```

```
print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)
```

```
Shape of the X_train (7000, 10)
Shape of the X_test (3000, 10)
Shape of the y_train (7000,)
Shape of the y_test (3000,)
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
# sequential model to initialise our ann and dense module to build the layers
from keras.models import Sequential
from keras.layers import Dense
```

```
2024-05-21 08:34:06.981674: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9373] Unable to register cuDNN factory: Attempt
2024-05-21 08:34:06.981725: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register cuFFT factory: Attempt
2024-05-21 08:34:06.982807: E external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1534] Unable to register cuBLAS factory: Attempt
2024-05-21 08:34:06.989509: I tensorflow/core/platform/cpu_feature_guard.cc:183] This TensorFlow binary is optimized to use availabl
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX, in other operations, rebuild TensorFlow with the appropriate compiler
```

```
classifier = Sequential()
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu', input_dim = 10))

# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))

# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
```

2024-05-21 08:34:10.706731: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1926] Created device /job:localhost/replica:0/task:0/

```
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

```
# Fitting the ANN to the Training set
```

```
classifier.fit(X_train, y_train, batch_size = 10, epochs = 100, verbose = 1)
```

```
Epoch 1/100
2024-05-21 08:34:30.051165: I external/local_xla/xla/service/service.cc:168] XLA service 0x7fc261d536a0 initialized for platform
2024-05-21 08:34:30.051210: I external/local_xla/xla/service/service.cc:176] StreamExecutor device (0): NVIDIA A100-SXM4-40GB M
2024-05-21 08:34:30.056612: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR crash reproducer, s
2024-05-21 08:34:30.108404: I external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:467] Loaded cuDNN version 90000
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1716280470.197288 7745 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for th
700/700 [=====] - 3s 1ms/step - loss: 0.4773 - accuracy: 0.7924
Epoch 2/100
700/700 [=====] - 1s 1ms/step - loss: 0.4296 - accuracy: 0.7926
Epoch 3/100
700/700 [=====] - 1s 1ms/step - loss: 0.4231 - accuracy: 0.8116
Epoch 4/100
700/700 [=====] - 1s 1ms/step - loss: 0.4188 - accuracy: 0.8240
Epoch 5/100
700/700 [=====] - 1s 1ms/step - loss: 0.4158 - accuracy: 0.8297
Epoch 6/100
700/700 [=====] - 1s 1ms/step - loss: 0.4134 - accuracy: 0.8314
Epoch 7/100
700/700 [=====] - 1s 1ms/step - loss: 0.4120 - accuracy: 0.8291
Epoch 8/100
700/700 [=====] - 1s 1ms/step - loss: 0.4107 - accuracy: 0.8336
Epoch 9/100
700/700 [=====] - 1s 1ms/step - loss: 0.4093 - accuracy: 0.8321
Epoch 10/100
700/700 [=====] - 1s 1ms/step - loss: 0.4089 - accuracy: 0.8327
Epoch 11/100
700/700 [=====] - 1s 1ms/step - loss: 0.4075 - accuracy: 0.8324
Epoch 12/100
700/700 [=====] - 1s 1ms/step - loss: 0.4071 - accuracy: 0.8320
Epoch 13/100
700/700 [=====] - 1s 1ms/step - loss: 0.4064 - accuracy: 0.8320
Epoch 14/100
700/700 [=====] - 1s 1ms/step - loss: 0.4050 - accuracy: 0.8336
Epoch 15/100
700/700 [=====] - 1s 1ms/step - loss: 0.4053 - accuracy: 0.8346
Epoch 16/100
700/700 [=====] - 1s 1ms/step - loss: 0.4045 - accuracy: 0.8326
Epoch 17/100
700/700 [=====] - 1s 1ms/step - loss: 0.4035 - accuracy: 0.8353
Epoch 18/100
700/700 [=====] - 1s 1ms/step - loss: 0.4041 - accuracy: 0.8359
Epoch 19/100
700/700 [=====] - 1s 1ms/step - loss: 0.4030 - accuracy: 0.8340
Epoch 20/100
700/700 [=====] - 1s 1ms/step - loss: 0.4020 - accuracy: 0.8344
Epoch 21/100
700/700 [=====] - 1s 1ms/step - loss: 0.4015 - accuracy: 0.8357
Epoch 22/100
700/700 [=====] - 1s 1ms/step - loss: 0.4015 - accuracy: 0.8331
Epoch 23/100
700/700 [=====] - 1s 1ms/step - loss: 0.4013 - accuracy: 0.8331
Epoch 24/100
700/700 [=====] - 1s 1ms/step - loss: 0.4010 - accuracy: 0.8356
Epoch 25/100
700/700 [=====] - 1s 1ms/step - loss: 0.3994 - accuracy: 0.8371
Epoch 26/100
```

```
score, acc = classifier.evaluate(X_train, y_train,
                                batch_size=10)
```

```
print('Train score:', score)
print('Train accuracy:', acc)
```

```
# Predicting the Test set results
```

```
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)
```

```
print('*'*20)
```

```
score, acc = classifier.evaluate(X_test, y_test,
                                batch_size=10)
```

```
print('Test score:', score)
print('Test accuracy:', acc)
```

```
700/700 [=====] - 1s 884us/step - loss: 0.3388 - accuracy: 0.8609
Train score: 0.33875057101249695
Train accuracy: 0.8608571290969849
94/94 [=====] - 0s 667us/step
```

300/300 [=====] - 0s 910us/step - loss: 0.3567 - accuracy: 0.8577
Test score: 0.3566587269306183
Test accuracy: 0.8576666712760925

```
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
target_names = ['Retained', 'Closed']
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

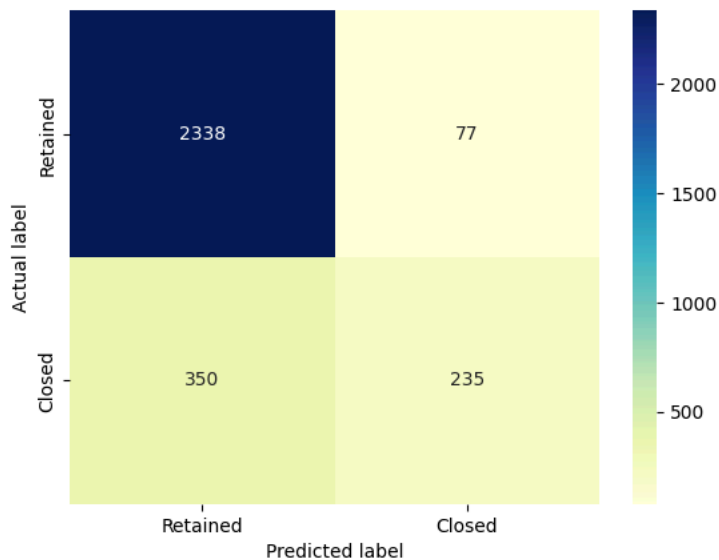
```
[[2338  77]
 [ 350 235]]
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names, yticklabels=target_names, cmap="YlGnBu", fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
Text(0.5, 23.52222222222222, 'Predicted label')
```

Confusion matrix



```
#import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred, target_names=target_names))
```


```
precision    recall  f1-score   support

Retained      0.87      0.97      0.92      2415
Closed        0.75      0.40      0.52       585

accuracy      0.81      0.68      0.72      3000
macro avg     0.81      0.68      0.72      3000
weighted avg  0.85      0.86      0.84      3000
```

```
from sklearn.metrics import roc_curve, auc
y_pred_proba = classifier.predict(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr, label='AUC (area = %0.2f)' % roc_auc)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.grid()
plt.legend(loc="lower right")
plt.title('ROC curve')
plt.show()
```

94/94 [=====] - 0s 667us/step



```
#Area under ROC curve
from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)
```

0.8357169400647662

✓ Experiment - 2

✓ ANN using Bird data

```
import numpy as np
import pandas as pd
```

```
bird_data = pd.read_csv('/workspace/FDP/Day3/bird.csv', delimiter = ',')
bird_data.head(5)
```

```
↗
```

	id	huml	humw	ulnal	ulnaw	feml	femw	tibl	tibw	tarl	tarw	type
0	0	80.78	6.68	72.01	4.88	41.81	3.70	5.50	4.03	38.70	3.84	SW
1	1	88.91	6.63	80.53	5.59	47.04	4.30	80.22	4.51	41.50	4.01	SW
2	2	79.97	6.37	69.26	5.28	43.07	3.90	75.35	4.04	38.31	3.34	SW
3	3	77.65	5.70	65.76	4.77	40.04	3.52	69.17	3.40	35.78	3.41	SW
4	4	62.80	4.84	52.09	3.73	33.95	2.72	56.27	2.96	31.88	3.13	SW

```
bird_data.columns
```

```
↗ Index(['id', 'huml', 'humw', 'ulnal', 'ulnaw', 'feml', 'femw', 'tibl', 'tibw',  
        'tarl', 'tarw', 'type'],  
        dtype='object')
```

```
bird_data = bird_data.set_index('id')
bird_data.head()
```

```
↗
```

	huml	humw	ulnal	ulnaw	feml	femw	tibl	tibw	tarl	tarw	type
id											
0	80.78	6.68	72.01	4.88	41.81	3.70	5.50	4.03	38.70	3.84	SW
1	88.91	6.63	80.53	5.59	47.04	4.30	80.22	4.51	41.50	4.01	SW
2	79.97	6.37	69.26	5.28	43.07	3.90	75.35	4.04	38.31	3.34	SW
3	77.65	5.70	65.76	4.77	40.04	3.52	69.17	3.40	35.78	3.41	SW
4	62.80	4.84	52.09	3.73	33.95	2.72	56.27	2.96	31.88	3.13	SW

```
bird_data.shape
```

```
↗ (420, 11)
```

```
bird_data.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
Int64Index: 420 entries, 0 to 419
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0    huml        419 non-null    float64
1    humw        419 non-null    float64
2    ulnal       417 non-null    float64
3    ulnaw       418 non-null    float64
4    feml        418 non-null    float64
5    femw        419 non-null    float64
6    tibl        418 non-null    float64
7    tibw        419 non-null    float64
8    tarl        419 non-null    float64
9    tarw        419 non-null    float64
10   type        420 non-null    object
dtypes: float64(10), object(1)
memory usage: 39.4+ KB
```

```
bird_data.isna().sum()
```

```
↗ huml      1
   humw      1
   ulnal      3
   ulnaw      2
   feml      2
   femw      1
```



```
tibl    2
tibw    1
tarl    1
tarw    1
type    0
dtype: int64
```

```
bird_data.dropna(how='any', inplace=True)
```

```
bird_data.isna().sum()
```

```
⇒ huml    0
   humw    0
   ulnal   0
   ulnaw   0
   feml    0
   femw    0
   tibl    0
   tibw    0
   tarl    0
   tarw    0
   type    0
dtype: int64
```

```
bird_data.shape
```

```
⇒ (413, 11)
```

```
bird_data.nunique()
```

```
⇒ huml    403
   humw   319
   ulnal   394
   ulnaw   305
   feml   397
   femw   287
   tibl   401
   tibw   283
   tarl   403
   tarw   277
   type     6
dtype: int64
```

```
bird_data['type'].unique()
```

```
⇒ array(['SW', 'W', 'T', 'R', 'P', 'SO'], dtype=object)
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
bird_data[['type']] = bird_data[['type']].apply(le.fit_transform)
```

```
bird_data.head()
```

```
⇒
```

	huml	humw	ulnal	ulnaw	feml	femw	tibl	tibw	tarl	tarw	type
id											
0	80.78	6.68	72.01	4.88	41.81	3.70	5.50	4.03	38.70	3.84	3
1	88.91	6.63	80.53	5.59	47.04	4.30	80.22	4.51	41.50	4.01	3
2	79.97	6.37	69.26	5.28	43.07	3.90	75.35	4.04	38.31	3.34	3
3	77.65	5.70	65.76	4.77	40.04	3.52	69.17	3.40	35.78	3.41	3
4	62.80	4.84	52.09	3.73	33.95	2.72	56.27	2.96	31.88	3.13	3

```
y = bird_data['type']
```

```
X = bird_data.drop(['type'],axis=1)
```

```
X.columns
```

```
⇒ Index(['huml', 'humw', 'ulnal', 'ulnaw', 'feml', 'femw', 'tibl', 'tibw',  
        'tarl', 'tarw'],  
       dtype='object')
```

```
y
```

```
⇒ id  
0    3
```

```

1      3
2      3
3      3
4      3
..
415    2
416    2
417    2
418    2
419    2
Name: type, Length: 413, dtype: int64

```

```
y.shape
```

```
(413,)
```

```

# from tensorflow.keras.utils import np_utils
# num_classes = 6
# y = np_utils.to_categorical(y, num_classes)
# y

```

```
from tensorflow.keras.utils import to_categorical
```

```

num_classes = 6
y = to_categorical(y, num_classes)
print(y)

```

```

[[0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 1. 0. 0.]
 ...
 [0. 0. 1. 0. 0. 0.]
 [0. 0. 1. 0. 0. 0.]
 [0. 0. 1. 0. 0. 0.]]

```

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2)

```

```

print("Shape of the X_train", X_train.shape)
print("Shape of the X_test", X_test.shape)
print("Shape of the y_train", y_train.shape)
print("Shape of the y_test", y_test.shape)

```

```

Shape of the X_train (330, 10)
Shape of the X_test (83, 10)
Shape of the y_train (330, 6)
Shape of the y_test (83, 6)

```

```

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

```

```

# sequential model to initialise our ann and dense module to build the layers
from keras.models import Sequential
from keras.layers import Dense

```

```

classifier = Sequential()
# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu', input_dim = 10))

```

```

# Adding the second hidden layer
classifier.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))

```

```

# Adding the third hidden layer
classifier.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu'))

```

```

# Adding the output layer
classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'softmax'))

```

```
2024-05-21 08:54:03.283163: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1926] Created device /job:localhost/replica:0/task:0/
```

```
classifier.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```

# Fitting the ANN to the Training set
classifier.fit(X_train, y_train, batch_size = 16, epochs = 800, verbose = 1)

```

```

Epoch 1/800
2024-05-21 08:54:07.013668: I external/local_xla/xla/service/service.cc:168] XLA service 0x7fd5fd374770 initialized for platform
2024-05-21 08:54:07.013710: I external/local_xla/xla/service/service.cc:176] StreamExecutor device (0): NVIDIA A100-SXM4-40GB M
2024-05-21 08:54:07.019479: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR crash reproducer, s
2024-05-21 08:54:07.057513: I external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:467] Loaded cuDNN version 90000
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1716281647.144364 13834 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for th
21/21 [=====] - 2s 2ms/step - loss: 1.7886 - accuracy: 0.2364
Epoch 2/800
21/21 [=====] - 0s 2ms/step - loss: 1.7772 - accuracy: 0.3364
Epoch 3/800
21/21 [=====] - 0s 2ms/step - loss: 1.7505 - accuracy: 0.4364
Epoch 4/800
21/21 [=====] - 0s 2ms/step - loss: 1.6777 - accuracy: 0.4818
Epoch 5/800
21/21 [=====] - 0s 2ms/step - loss: 1.5472 - accuracy: 0.4909
Epoch 6/800
21/21 [=====] - 0s 2ms/step - loss: 1.4310 - accuracy: 0.4970
Epoch 7/800
21/21 [=====] - 0s 2ms/step - loss: 1.3724 - accuracy: 0.5000
Epoch 8/800
21/21 [=====] - 0s 2ms/step - loss: 1.3375 - accuracy: 0.5030
Epoch 9/800
21/21 [=====] - 0s 2ms/step - loss: 1.3134 - accuracy: 0.5030
Epoch 10/800
21/21 [=====] - 0s 2ms/step - loss: 1.2927 - accuracy: 0.5061
Epoch 11/800
21/21 [=====] - 0s 2ms/step - loss: 1.2780 - accuracy: 0.5091
Epoch 12/800
21/21 [=====] - 0s 2ms/step - loss: 1.2687 - accuracy: 0.5091
Epoch 13/800
21/21 [=====] - 0s 2ms/step - loss: 1.2571 - accuracy: 0.5121
Epoch 14/800
21/21 [=====] - 0s 2ms/step - loss: 1.2445 - accuracy: 0.5212
Epoch 15/800
21/21 [=====] - 0s 2ms/step - loss: 1.2328 - accuracy: 0.5242
Epoch 16/800
21/21 [=====] - 0s 2ms/step - loss: 1.2212 - accuracy: 0.5273
Epoch 17/800
21/21 [=====] - 0s 2ms/step - loss: 1.2172 - accuracy: 0.5273
Epoch 18/800
21/21 [=====] - 0s 2ms/step - loss: 1.2037 - accuracy: 0.5303
Epoch 19/800
21/21 [=====] - 0s 2ms/step - loss: 1.1877 - accuracy: 0.5303
Epoch 20/800
21/21 [=====] - 0s 2ms/step - loss: 1.1763 - accuracy: 0.5333
Epoch 21/800
21/21 [=====] - 0s 2ms/step - loss: 1.1673 - accuracy: 0.5364
Epoch 22/800
21/21 [=====] - 0s 2ms/step - loss: 1.1575 - accuracy: 0.5394
Epoch 23/800
21/21 [=====] - 0s 2ms/step - loss: 1.1437 - accuracy: 0.5424
Epoch 24/800
21/21 [=====] - 0s 2ms/step - loss: 1.1276 - accuracy: 0.5424
Epoch 25/800
21/21 [=====] - 0s 2ms/step - loss: 1.1143 - accuracy: 0.5455
Epoch 26/800

```

```

score, acc = classifier.evaluate(X_train, y_train,
                                batch_size=10)
print('Train score:', score)
print('Train accuracy:', acc)

print('*'*20)
score, acc = classifier.evaluate(X_test, y_test,
                                batch_size=10)
print('Test score:', score)
print('Test accuracy:', acc)

33/33 [=====] - 0s 999us/step - loss: 0.1421 - accuracy: 0.9545
Train score: 0.14205241203308105
Train accuracy: 0.9545454382896423
*****
9/9 [=====] - 0s 1ms/step - loss: 0.4386 - accuracy: 0.8916
Test score: 0.4385785162448883
Test accuracy: 0.891566276550293

# Predicting the Test set results
pred = classifier.predict(X_test)
print("Y_pred:", pred)
print("*****")
y_pred = np.argmax(pred, axis = 1)
print("Y_pred:", y_pred)
print("*****")
print("Y_test:", y_test)
y_true = np.argmax(y_test, axis = 1)

```


[illegible]

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
p = sns.heatmap(pd.DataFrame(cm), annot=True, xticklabels=target_names, yticklabels=target_names, cmap="YlGnBu", fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

Text(0.5, 23.52222222222222, 'Predicted label')

Confusion matrix



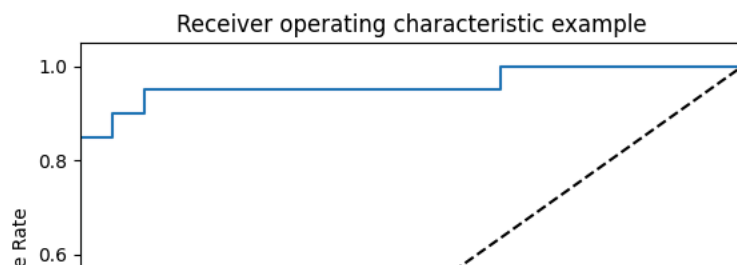



```
#import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_true,y_pred, target_names = target_names))
```




	precision	recall	f1-score	support
P	0.80	0.67	0.73	6
R	0.91	0.91	0.91	11
SO	0.89	1.00	0.94	31
SW	0.94	0.85	0.89	20
T	0.75	1.00	0.86	3
W	0.90	0.75	0.82	12
accuracy			0.89	83
macro avg	0.86	0.86	0.86	83
weighted avg	0.89	0.89	0.89	83

✓ ROC curve

```
from sklearn.metrics import roc_curve, auc
from itertools import cycle
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(6):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot of a ROC curve for a specific class
for i in range(6):
    plt.figure()
    plt.plot(fpr[i], tpr[i], label='ROC curve (area = %0.2f)' % roc_auc[i])
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic example')
    plt.legend(loc="lower right")
    plt.show()
```





```

fpr = dict()
tpr = dict()
roc_auc = dict()
lw=2
for i in range(6):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
colors =cycle(['blue', 'green', 'red', 'darkorange', 'olive', 'purple'])
for i, color in zip(range(6), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='AUC = {1:0.4f}'
             ''.format(i, roc_auc[i]))
plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=15)
plt.ylabel('True Positive Rate', fontsize=15)
# plt.title('Receiver operating characteristic for multi-class data')
plt.legend(loc="lower right")
plt.show()

```



✓ Experiment - 3

```
from tensorflow.keras import applications
from tensorflow.keras import optimizers
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from tensorflow.keras import backend as k
from tensorflow.keras.callbacks import ModelCheckpoint, LearningRateScheduler, TensorBoard, EarlyStopping

import numpy as np
from tensorflow.keras import models
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model
from tensorflow.keras import layers
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras import Input

# Normalize training and validation data in the range of 0 to 1
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
                                                    # horizontal_flip=True,
                                                    # height_shift_range=0.1,
                                                    # width_shift_range=0.1
validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)


# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    'C://Users//abhia//Downloads//plant_village (1)//plant_village/t/rain',
    target_size=(64, 64),
    batch_size=16,
    class_mode='categorical')

# Read Validation data from directory and define target size with batch size
validation_generator = validation_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/val/',
    target_size=(64, 64),
    batch_size=16,
    class_mode='categorical',
    shuffle=False)

test_generator = test_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/test/',
    target_size=(64, 64),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

Visualization of few images

```
plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
img, label = test_generator.next()
img[0].shape
```

```
(64, 64, 3)
```

```
# Create the model
model = models.Sequential()
# Add new layers
model.add(Conv2D(128, kernel_size=(3,3), activation = 'relu', input_shape=(64,64,3)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(64, kernel_size=(3,3), activation = 'relu'))
```

```

model.add(MaxPooling2D(pool_size=(2,2)))
model.add(layers.Flatten())
model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dense(4, activation='softmax'))
model.summary()

```

Compiling and Training the Model

```

# We are going to use accuracy metrics and cross entropy loss as performance parameters
model.compile(optimizer = optimizers.Adam(learning_rate = 0.0001), loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=30,
                    validation_data=validation_generator,
                    validation_steps=validation_generator.samples/validation_generator.batch_size,
                    verbose=2)

```



```

Epoch 1/30
Epoch 1/30
188/187 - 2s - loss: 0.0354 - acc: 0.9893 - val_loss: 0.2115 - val_acc: 0.9402
Epoch 2/30
Epoch 1/30
188/187 - 2s - loss: 0.0313 - acc: 0.9890 - val_loss: 0.2183 - val_acc: 0.9339
Epoch 3/30
Epoch 1/30
188/187 - 2s - loss: 0.0292 - acc: 0.9903 - val_loss: 0.2038 - val_acc: 0.9465
Epoch 4/30
Epoch 1/30
188/187 - 2s - loss: 0.0314 - acc: 0.9907 - val_loss: 0.2843 - val_acc: 0.9181
Epoch 5/30
Epoch 1/30
188/187 - 2s - loss: 0.0213 - acc: 0.9940 - val_loss: 0.3160 - val_acc: 0.9181
Epoch 6/30
Epoch 1/30
188/187 - 2s - loss: 0.0276 - acc: 0.9923 - val_loss: 0.2019 - val_acc: 0.9449
Epoch 7/30
Epoch 1/30
188/187 - 2s - loss: 0.0184 - acc: 0.9960 - val_loss: 0.2326 - val_acc: 0.9307
Epoch 8/30
Epoch 1/30
188/187 - 2s - loss: 0.0161 - acc: 0.9970 - val_loss: 0.2079 - val_acc: 0.9386
Epoch 9/30
Epoch 1/30
188/187 - 2s - loss: 0.0145 - acc: 0.9980 - val_loss: 0.2025 - val_acc: 0.9465
Epoch 10/30
Epoch 1/30
188/187 - 2s - loss: 0.0180 - acc: 0.9953 - val_loss: 0.2236 - val_acc: 0.9386
Epoch 11/30
Epoch 1/30
188/187 - 2s - loss: 0.0152 - acc: 0.9970 - val_loss: 0.2413 - val_acc: 0.9339
Epoch 12/30
Epoch 1/30
188/187 - 2s - loss: 0.0245 - acc: 0.9913 - val_loss: 0.1965 - val_acc: 0.9465
Epoch 13/30
Epoch 1/30
188/187 - 2s - loss: 0.0202 - acc: 0.9940 - val_loss: 0.3189 - val_acc: 0.9228
Epoch 14/30
Epoch 1/30
188/187 - 2s - loss: 0.0136 - acc: 0.9970 - val_loss: 0.1991 - val_acc: 0.9449
Epoch 15/30
Epoch 1/30
188/187 - 2s - loss: 0.0083 - acc: 0.9983 - val_loss: 0.2098 - val_acc: 0.9402
Epoch 16/30
Epoch 1/30
188/187 - 2s - loss: 0.0081 - acc: 0.9993 - val_loss: 0.2170 - val_acc: 0.9496
Epoch 17/30
Epoch 1/30
188/187 - 2s - loss: 0.0104 - acc: 0.9980 - val_loss: 0.2084 - val_acc: 0.9480
Epoch 18/30
Epoch 1/30
188/187 - 2s - loss: 0.0070 - acc: 0.9993 - val_loss: 0.1953 - val_acc: 0.9480
Epoch 19/30
Epoch 1/30
188/187 - 2s - loss: 0.0578 - acc: 0.9804 - val_loss: 0.4641 - val_acc: 0.8898
Epoch 20/30

```

Saving the model

```

model.save("CONV_plant_deseas.h5")
print("Saved model to disk")

```



Saved model to disk

Loading the model

```
model = models.load_model('CONV_plant_deseas.h5')
```

Visualization of Accuracy and Loss Curves

```
train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```

Prediction

```
# Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size, verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}".format(len(errors),test_generator.samples))
```


```
546/546 [=====] - 1s 1ms/step
No of errors = 38/546
```

```
accuracy = ((test_generator.samples-len(errors))/test_generator.samples) * 100
accuracy
```

```
93.04029304029304
```

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names=label2index))
```

	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.91	0.92	0.92	103
Apple__Apple_scab	0.90	0.98	0.94	134
Apple__Cedar_apple_rust	0.85	0.71	0.78	49
Apple__healthy	0.97	0.95	0.96	260
accuracy			0.93	546
macro avg	0.91	0.89	0.90	546
weighted avg	0.93	0.93	0.93	546

Experiment - 4

✓ Deep Learning Training and Architecture, Feature Extraction, Models training with some pretrained models.

```
import numpy as np
from tensorflow.keras import Input
from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
from tensorflow.keras.models import Model
from tensorflow.keras import applications
from tensorflow.keras import backend as k
import matplotlib.pyplot as plt
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from tensorflow.keras.callbacks import ModelCheckpoint, LearningRateScheduler, TensorBoard, EarlyStopping
```

```
# Normalize training and validation data in the range of 0 to 1
train_datagen = ImageDataGenerator(rescale=1./255) # vertical_flip=True,
                                                    # horizontal_flip=True,
                                                    # height_shift_range=0.1,
                                                    # width_shift_range=0.1

validation_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)
```


```
# Read the training sample and set the batch size
train_generator = train_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/train/',
    target_size=(128, 128),
    batch_size=16,
    class_mode='categorical')
```

```
# Read Validation data from directory and define target size with batch size
validation_generator = validation_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/val/',
    target_size=(128, 128),
    batch_size=16,
    class_mode='categorical',
    shuffle=False)
```

```
test_generator = test_datagen.flow_from_directory(
    '/workspace/Bootcamp/Data/plant_village/test/',
    target_size=(128, 128),
    batch_size=1,
    class_mode='categorical',
    shuffle=False)
```

🔍 Found 3004 images belonging to 4 classes.
Found 635 images belonging to 4 classes.
Found 547 images belonging to 4 classes.

```
plt.figure(figsize=(16, 16))
for i in range(1, 17):
    plt.subplot(4, 4, i)
    img, label = test_generator.next()
    # print(img.shape)
    # print(label)
    plt.imshow(img[0])
plt.show()
```



```
img, label = test_generator.next()
img[0].shape
```

(128, 128, 3)

✓ VGG16

```
from tensorflow.keras.applications.vgg16 import VGG16
```

```
## Loading VGG16 model
```

```
base_model = VGG16(weights="imagenet", include_top=False, input_shape= (128, 128, 3))
# Include_top = False means excluding the model fully connected layers
base_model.trainable = False ## Not trainable weights,
#weights of the VGG16 model will not be updated during training
base_model.summary()
```



Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 128, 128, 3)]	0

block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792

block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928

block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0

block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856

block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584

block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0

block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168

block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080

block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080

block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0

block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160

block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808

block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808

block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0

block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808

block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808

block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808

block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(64, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    prediction_layer
])
model.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Model)	(None, 4, 4, 512)	14714688

global_average_pooling2d (Gl	(None, 512)	0

dense (Dense)	(None, 4)	2052
=====		
Total params: 14,716,740		
Trainable params: 2,052		
Non-trainable params: 14,714,688		

✓ Training

```
# sgd = SGD(lr=0.001,decay=1e-6, momentum=0.9, nesterov=True)
# We are going to use accuracy metrics and cross entropy loss as performance parameters
```



```

model.compile(optimizer = Adam(learning_rate = 0.0001), loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=30,
                    validation_data=validation_generator,
                    validation_steps=validation_generator.samples/validation_generator.batch_size,
                    verbose=1)

```

188/187 [=====] - 6s 32ms/step - loss: 0.1155 - acc: 0.9710 - val_loss: 0.1342 - val_acc: 0.9591
 Epoch 12/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1159 - acc: 0.9701Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1150 - acc: 0.9704 - val_loss: 0.1345 - val_acc: 0.9559
 Epoch 13/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1138 - acc: 0.9707Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1149 - acc: 0.9704 - val_loss: 0.1343 - val_acc: 0.9559
 Epoch 14/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1142 - acc: 0.9717Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1145 - acc: 0.9710 - val_loss: 0.1334 - val_acc: 0.9559
 Epoch 15/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1147 - acc: 0.9707Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1143 - acc: 0.9710 - val_loss: 0.1335 - val_acc: 0.9591
 Epoch 16/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1138 - acc: 0.9714Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1139 - acc: 0.9714 - val_loss: 0.1339 - val_acc: 0.9575
 Epoch 17/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1135 - acc: 0.9717Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1138 - acc: 0.9714 - val_loss: 0.1343 - val_acc: 0.9575
 Epoch 18/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1135 - acc: 0.9704Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1134 - acc: 0.9704 - val_loss: 0.1340 - val_acc: 0.9591
 Epoch 19/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1134 - acc: 0.9714Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1137 - acc: 0.9714 - val_loss: 0.1339 - val_acc: 0.9559
 Epoch 20/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1134 - acc: 0.9717Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1130 - acc: 0.9717 - val_loss: 0.1322 - val_acc: 0.9575
 Epoch 21/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1124 - acc: 0.9711Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1126 - acc: 0.9710 - val_loss: 0.1325 - val_acc: 0.9559
 Epoch 22/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1127 - acc: 0.9714Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1123 - acc: 0.9717 - val_loss: 0.1333 - val_acc: 0.9591
 Epoch 23/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1108 - acc: 0.9714Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1120 - acc: 0.9707 - val_loss: 0.1331 - val_acc: 0.9591
 Epoch 24/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1117 - acc: 0.9717Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1117 - acc: 0.9714 - val_loss: 0.1324 - val_acc: 0.9559
 Epoch 25/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1111 - acc: 0.9711Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1114 - acc: 0.9710 - val_loss: 0.1318 - val_acc: 0.9559
 Epoch 26/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1114 - acc: 0.9707Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1113 - acc: 0.9710 - val_loss: 0.1314 - val_acc: 0.9559
 Epoch 27/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1093 - acc: 0.9727Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1109 - acc: 0.9717 - val_loss: 0.1318 - val_acc: 0.9559
 Epoch 28/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1108 - acc: 0.9721Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1106 - acc: 0.9720 - val_loss: 0.1307 - val_acc: 0.9575
 Epoch 29/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1104 - acc: 0.9711Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1104 - acc: 0.9710 - val_loss: 0.1322 - val_acc: 0.9575
 Epoch 30/30
 186/187 [=====>.] - ETA: 0s - loss: 0.1096 - acc: 0.9738Epoch 1/30
 188/187 [=====] - 6s 32ms/step - loss: 0.1101 - acc: 0.9734 - val_loss: 0.1305 - val_acc: 0.9559


```

model.save("VGG16_plant_deseas.h5")
print("Saved model to disk")

```

Saved model to disk

```

model = models.load_model('VGG16_plant_deseas.h5')
print("Model is loaded")

```

Model is loaded

```

model.save_weights('cnn_classification.h5')

model.load_weights('cnn_classification.h5')


train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']

```


```
val_loss = history.history['val_loss']
```

```
epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()
```

```
plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>



```
# Get the filenames from the generator
fnames = test_generator.filenames
```

```
# Get the ground truth from generator
ground_truth = test_generator.classes
```

```
# Get the label to class mapping from the generator
label2index = test_generator.class_indices
```

```
# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())
```

```
# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)
```

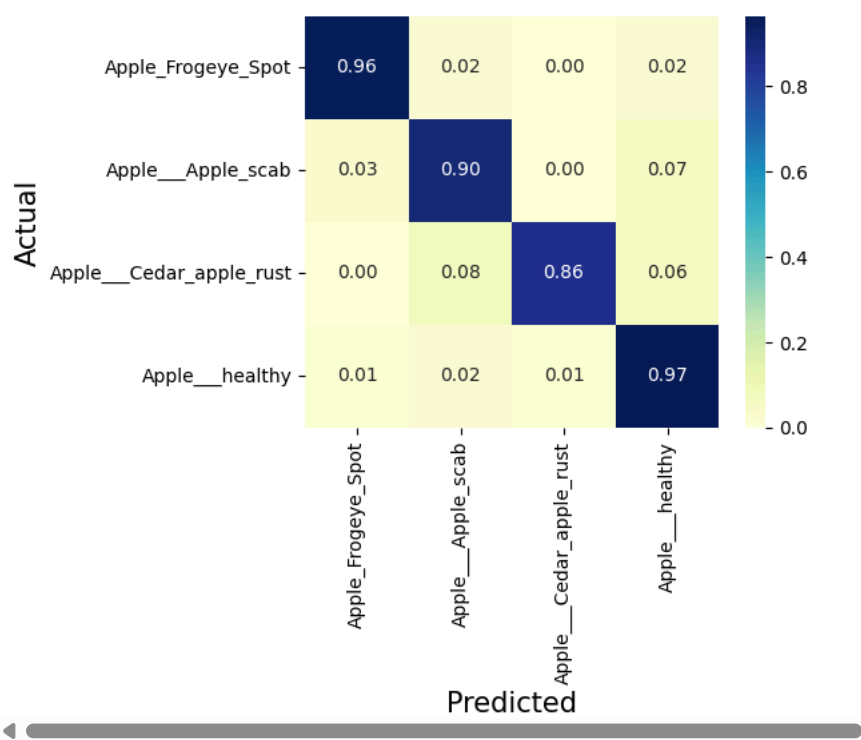
```
errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}".format(len(errors),test_generator.samples))
```

```
547/547 [=====] - 3s 6ms/step
No of errors = 33/547
```

```
accuracy = ((test_generator.samples-len(errors))/test_generator.samples) * 100
accuracy
```

```
93.96709323583181
```

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)
```



```
from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names=label2index))
```

```
precision    recall  f1-score   support

Apple_Frogeye_Spot      0.94      0.96      0.95        104
Apple__Apple_scab       0.92      0.90      0.91        134
Apple__Cedar_apple_rust  0.95      0.86      0.90         49
Apple__healthy          0.95      0.97      0.96        260

   accuracy                   0.94        547
  macro avg       0.94      0.92      0.93        547
 weighted avg     0.94      0.94      0.94        547
```

Start coding or generate with AI.

```
from tensorflow.keras import applications

## Loading InceptionV3 model
base_model = applications.InceptionV3(weights="imagenet", include_top=False, input_shape= (128, 128, 3))
base_model.trainable = False ## Not trainable weights

base_model.summary()
```

base_model.summary()

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_orderin87916544/87910968 [=====] - 15s 0us/step
Model: "inception_v3"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_3 (InputLayer)	(None, 128, 128, 3)	0	
conv2d (Conv2D)	(None, 63, 63, 32)	864	input_3[0][0]
batch_normalization (BatchNormaliza	(None, 63, 63, 32)	96	conv2d[0][0]
activation (Activation)	(None, 63, 63, 32)	0	batch_normalization[0][0]
conv2d_1 (Conv2D)	(None, 61, 61, 32)	9216	activation[0][0]
batch_normalization_1 (BatchNor	(None, 61, 61, 32)	96	conv2d_1[0][0]
activation_1 (Activation)	(None, 61, 61, 32)	0	batch_normalization_1[0][0]
conv2d_2 (Conv2D)	(None, 61, 61, 64)	18432	activation_1[0][0]
batch_normalization_2 (BatchNor	(None, 61, 61, 64)	192	conv2d_2[0][0]
activation_2 (Activation)	(None, 61, 61, 64)	0	batch_normalization_2[0][0]
max_pooling2d (MaxPooling2D)	(None, 30, 30, 64)	0	activation_2[0][0]
conv2d_3 (Conv2D)	(None, 30, 30, 80)	5120	max_pooling2d[0][0]
batch_normalization_3 (BatchNor	(None, 30, 30, 80)	240	conv2d_3[0][0]
activation_3 (Activation)	(None, 30, 30, 80)	0	batch_normalization_3[0][0]
conv2d_4 (Conv2D)	(None, 28, 28, 192)	138240	activation_3[0][0]
batch_normalization_4 (BatchNor	(None, 28, 28, 192)	576	conv2d_4[0][0]
activation_4 (Activation)	(None, 28, 28, 192)	0	batch_normalization_4[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 192)	0	activation_4[0][0]
conv2d_8 (Conv2D)	(None, 13, 13, 64)	12288	max_pooling2d_1[0][0]
batch_normalization_8 (BatchNor	(None, 13, 13, 64)	192	conv2d_8[0][0]
activation_8 (Activation)	(None, 13, 13, 64)	0	batch_normalization_8[0][0]
conv2d_6 (Conv2D)	(None, 13, 13, 48)	9216	max_pooling2d_1[0][0]
conv2d_9 (Conv2D)	(None, 13, 13, 96)	55296	activation_8[0][0]
batch_normalization_6 (BatchNor	(None, 13, 13, 48)	144	conv2d_6[0][0]
batch_normalization_9 (BatchNor	(None, 13, 13, 96)	288	conv2d_9[0][0]
activation_6 (Activation)	(None, 13, 13, 48)	0	batch_normalization_6[0][0]

```
flatten_layer = layers.GlobalAveragePooling2D()
dense_layer_1 = layers.Dense(64, activation='relu')
dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')
```

```
model = models.Sequential([
    base_model,
    flatten_layer,
    dense_layer_1,
    dense_layer_2,
    prediction_layer
])
```

model.summary()

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
inception_v3 (Model)	(None, 2, 2, 2048)	21802784
global_average_pooling2d_2 ((None, 2048)	0
dense_2 (Dense)	(None, 64)	131136
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 4)	132

```
=====
Total params: 21,936,132
Trainable params: 133,348
Non-trainable params: 21,802,784
=====
```

```
model.compile(optimizer = Adam(learning_rate = 0.001), loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=10,
                    validation_data=validation_generator,
                    validation_steps=validation_generator.samples/validation_generator.batch_size,
                    verbose=2)
```

```
↔ Epoch 1/10
Epoch 1/10
188/187 - 9s - loss: 0.3544 - acc: 0.8725 - val_loss: 1.8910 - val_acc: 0.7134
Epoch 2/10
Epoch 1/10
188/187 - 5s - loss: 0.4988 - acc: 0.8113 - val_loss: 1.5254 - val_acc: 0.7181
Epoch 3/10
Epoch 1/10
188/187 - 5s - loss: 0.4549 - acc: 0.8246 - val_loss: 1.2376 - val_acc: 0.7433
Epoch 4/10
Epoch 1/10
188/187 - 5s - loss: 0.4242 - acc: 0.8342 - val_loss: 1.2466 - val_acc: 0.7780
Epoch 5/10
Epoch 1/10
188/187 - 5s - loss: 0.4462 - acc: 0.8382 - val_loss: 1.6309 - val_acc: 0.7165
Epoch 6/10
Epoch 1/10
188/187 - 5s - loss: 0.4180 - acc: 0.8445 - val_loss: 1.3079 - val_acc: 0.7260
Epoch 7/10
Epoch 1/10
188/187 - 5s - loss: 0.4094 - acc: 0.8442 - val_loss: 2.0786 - val_acc: 0.6693
Epoch 8/10
Epoch 1/10
188/187 - 5s - loss: 0.3816 - acc: 0.8595 - val_loss: 1.5253 - val_acc: 0.7606
Epoch 9/10
Epoch 1/10
188/187 - 5s - loss: 0.4081 - acc: 0.8449 - val_loss: 1.1693 - val_acc: 0.7606
Epoch 10/10
Epoch 1/10
188/187 - 5s - loss: 0.3675 - acc: 0.8615 - val_loss: 1.4619 - val_acc: 0.7276
```


```
model.save("InceptionNet_plant_deseas.h5")
print("Saved model to disk")
```

```
↔ Saved model to disk
```


```
train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()
```

```
plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()
```



<Figure size 640x480 with 0 Axes>



```
# Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}".format(len(errors),test_generator.samples))
```



547/547 [=====] - 9s 17ms/step
No of errors = 153/547

```
accuracy = ((test_generator.samples-len(errors))/test_generator.samples) * 100
accuracy
```



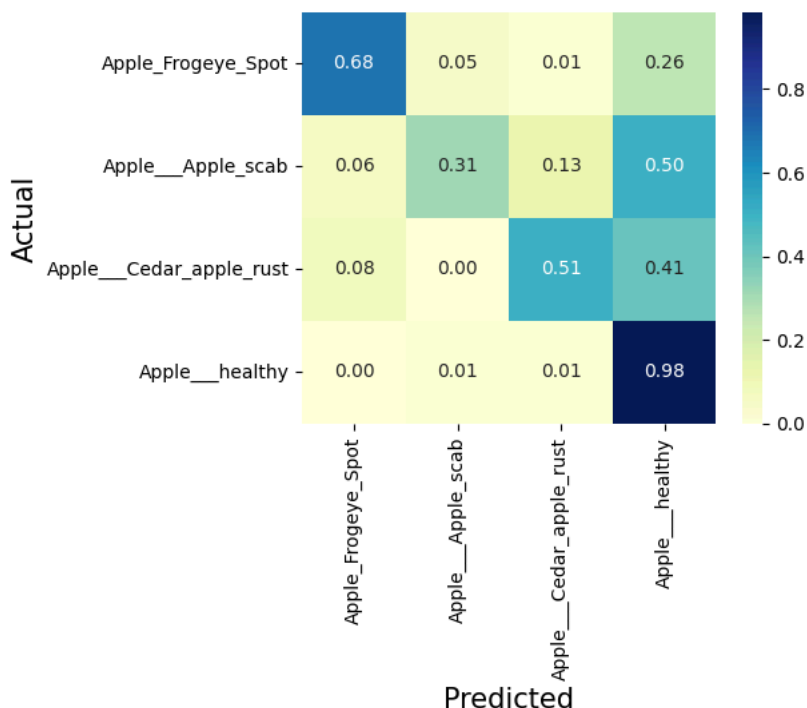
72.0292504570384

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
```

```

cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)

```



```

from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names=label2index))

```



	precision	recall	f1-score	support
Apple_Frogeye_Spot	0.86	0.68	0.76	104
Apple__Apple_scab	0.86	0.31	0.46	134
Apple__Cedar_apple_rust	0.56	0.51	0.53	49
Apple__healthy	0.69	0.98	0.81	260
accuracy			0.72	547
macro avg	0.74	0.62	0.64	547
weighted avg	0.75	0.72	0.69	547

Start coding or generate with AI.

▼ ResNet

```

from keras import applications

```

```

## Loading VGG16 model
base_model = applications.ResNet50(weights="imagenet", include_top=False, input_shape= (128, 128, 3))
base_model.trainable = False ## Not trainable weights

```

```

base_model.summary()

```

```

flatten_layer = layers.GlobalAveragePooling2D()
# dense_layer_1 = layers.Dense(63, activation='relu')
# dense_layer_2 = layers.Dense(32, activation='relu')
prediction_layer = layers.Dense(4, activation='softmax')

```

```

model = models.Sequential([
    base_model,
    flatten_layer,
    # dense_layer_1,
    # dense_layer_2,
    prediction_layer
])

```

```

model.summary()

model.compile(optimizer = Adam(learning_rate = 0.001), loss='categorical_crossentropy', metrics=['acc'])
# Train the model
history = model.fit(train_generator,
                    steps_per_epoch=train_generator.samples/train_generator.batch_size,
                    epochs=30,
                    validation_data=validation_generator,
                    validation_steps=validation_generator.samples/validation_generator.batch_size,
                    verbose=1)

model.save("ResNet_plant_deseas.h5")
print("Saved model to disk")

train_acc = history.history['acc']
val_acc = history.history['val_acc']
train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(train_acc))
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'g', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.grid()
plt.legend()
plt.figure()
plt.show()

plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'g', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.grid()
plt.legend()
plt.show()

# Get the filenames from the generator
fnames = test_generator.filenames

# Get the ground truth from generator
ground_truth = test_generator.classes

# Get the label to class mapping from the generator
label2index = test_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(test_generator, steps=test_generator.samples/test_generator.batch_size, verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}".format(len(errors),test_generator.samples))

accuracy = ((test_generator.samples-len(errors))/test_generator.samples) * 100
accuracy

from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
cm = confusion_matrix(y_true=ground_truth, y_pred=predicted_classes)
cm = np.array(cm)
# Normalise
cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
fig, ax = plt.subplots(figsize=(5,4))
sns.heatmap(cmn, annot=True, fmt='.2f', xticklabels=label2index, yticklabels=label2index, cmap="YlGnBu")
plt.ylabel('Actual', fontsize=15)
plt.xlabel('Predicted', fontsize=15)
plt.show(block=False)

from sklearn.metrics import classification_report
print(classification_report(ground_truth, predicted_classes, target_names=label2index))

```


Experiment - 5

✓ Text data handling with RNN for sentiment analysis

```
import pandas as pd      # to load dataset
import numpy as np       # for mathematic equation
from nltk.corpus import stopwords  # to get collection of stopwords
from sklearn.model_selection import train_test_split  # for splitting dataset
from tensorflow.keras.preprocessing.text import Tokenizer  # to encode text to int
from tensorflow.keras.preprocessing.sequence import pad_sequences  # to do padding or truncating
from tensorflow.keras.models import Sequential  # the model
from tensorflow.keras.layers import Embedding, LSTM, Dense  # layers of the architecture
from tensorflow.keras.callbacks import ModelCheckpoint  # save model
from tensorflow.keras.models import load_model  # load saved model
import re
from keras.layers import SimpleRNN
```

```
data = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/IMDB Dataset.csv')
print(data)
```

```

review sentiment
0      One of the other reviewers has mentioned that ... positive
1      A wonderful little production. <br /><br />The... positive
2      I thought this was a wonderful way to spend ti... positive
3      Basically there's a family where a little boy ... negative
4      Petter Mattei's "Love in the Time of Money" is... positive
...
49995  I thought this movie did a down right good job... positive
49996  Bad plot, bad dialogue, bad acting, idiotic di... negative
49997  I am a Catholic taught in parochial elementary... negative
49998  I'm going to have to disagree with the previou... negative
49999  No one expects the Star Trek movies to be high... negative

[50000 rows x 2 columns]
```

Stop Word is a commonly used words in a sentence, usually a search engine is programmed to ignore this words (i.e. "the", "a", "an", "of", etc.)

Declaring the english stop words

```
import nltk
nltk.download("stopwords")
english_stops = set(stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
def load_dataset():
    df = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/IMDB Dataset.csv')
    x_data = df['review']  # Reviews/Input
    y_data = df['sentiment']  # Sentiment/Output

    # PRE-PROCESS REVIEW
    x_data = x_data.replace({'<.*?>': ''}, regex = True)  # remove html tag
    x_data = x_data.replace({'^[A-Za-z]': ' '}, regex = True)  # remove non alphabet
    x_data = x_data.apply(lambda review: [w for w in review.split() if w not in english_stops])  # remove stop words
    x_data = x_data.apply(lambda review: [w.lower() for w in review])  # lower case

    # ENCODE SENTIMENT -> 0 & 1
    y_data = y_data.replace('positive', 1)
    y_data = y_data.replace('negative', 0)

    return x_data, y_data

x_data, y_data = load_dataset()

print('Reviews')
print(x_data, '\n')
print('Sentiment')
print(y_data)
```

```

Reviews
0      [one, reviewers, mentioned, watching, oz, epis...
1      [a, wonderful, little, production, the, filmin...
2      [i, thought, wonderful, way, spend, time, hot,...
3      [basically, family, little, boy, jake, thinks,...
4      [petter, mattei, love, time, money, visually, ...
```

```

...
49995 [i, thought, movie, right, good, job, it, crea...
49996 [bad, plot, bad, dialogue, bad, acting, idioti...
49997 [i, catholic, taught, parochial, elementary, s...
49998 [i, going, disagree, previous, comment, side, ...
49999 [no, one, expects, star, trek, movies, high, a...
Name: review, Length: 50000, dtype: object

```

```

Sentiment
0      1
1      1
2      1
3      0
4      1
..
49995   1
49996   0
49997   0
49998   0
49999   0
Name: sentiment, Length: 50000, dtype: int64

```

```
x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size = 0.2)
```

```

print('Train Set')
print(x_train, '\n')
print(x_test, '\n')
print('Test Set')
print(y_train, '\n')
print(y_test)

```

```

➡ Train Set
22596 [boring, badly, written, italian, exploitation...
5353  [the, other, supposed, horror, movie, made, it...
42152 [a, tough, life, gets, tougher, three, childre...
15434 [why, earth, colin, firth, pointless, film, ha...
7280  [this, far, worst, movie, i, ever, seen, cinem...
...
39945 [this, show, lasted, moments, plots, usually, ...
13858 [i, rented, thinking, would, pretty, good, cov...
25266 [having, pleasantly, surprised, sandra, bulloc...
10659 [the, difficulty, i, musical, version, les, mi...
39372 [this, movie, proof, film, noire, enduring, st...
Name: review, Length: 40000, dtype: object

2006  [this, movie, time, favorite, you, really, see...
33575 [this, british, film, version, stage, play, i,...
6808  [alexander, nevsky, brilliant, piece, cinemati...
32330 [found, old, vhs, version, film, parents, hous...
3777  [i, went, see, movie, daughter, i, insisted, g...
...
40255 [what, heck, people, expect, horror, films, da...
5864  [especially, time, much, science, fiction, fil...
44604 [nicole, eggert, listed, star, despite, michea...
42481 [a, thief, night, got, best, end, times, thril...
31671 [i, enjoy, national, anthem, i, enjoy, nationa...
Name: review, Length: 10000, dtype: object

```

```

Test Set
22596   0
5353    0
42152   1
15434   0
7280    0
..
39945   0
13858   0
25266   0
10659   0
39372   1
Name: sentiment, Length: 40000, dtype: int64

2006    1
33575   1
6808    1
32330   0
3777    0
..
40255   1
5864    1
44604   0
42481   1
31671   1
Name: sentiment, Length: 10000, dtype: int64

```

Function for getting the average review length, by calculating the mean of all the reviews length (using numpy.mean)

```

def get_max_length():
    review_length = []
    for review in x_train:
        review_length.append(len(review))

    return int(np.ceil(np.mean(review_length)))

# ENCODE REVIEW
token = Tokenizer(lower=False) # no need lower, because already lowered the data in load_data()
token.fit_on_texts(x_train)
x_train = token.texts_to_sequences(x_train)
x_test = token.texts_to_sequences(x_test)

max_length = get_max_length()

x_train = pad_sequences(x_train, maxlen=max_length, padding='post', truncating='post')
x_test = pad_sequences(x_test, maxlen=max_length, padding='post', truncating='post')

total_words = len(token.word_index) + 1 # add 1 because of 0 padding
print('Total Words:', total_words)

print('Encoded X Train\n', x_train, '\n')
print('Encoded X Test\n', x_test, '\n')
print('Maximum review length: ', max_length)

```

```

→ Total Words: 92636
Encoded X Train
[[ 257  863  310 ...    0    0    0]
 [   2 1340  350 ...   28  282 409]
 [  39 1138   40 ...    0    0    0]
 ...
 [1587 3903  660 ...   62 14457 1006]
 [   2 6090   1 ...  4973  5675  406]
 [   8   3 2912 ...    0    0    0]]

Encoded X Test
[[   8   3  10 ...    0    0    0]
 [   8  603   4 ...  278 10278 2289]
 [3551 11276  417 ...    0    0    0]
 ...
 [3908 20405 3718 ...    0    0    0]
 [  39 2984  218 ...  3947   3  765]
 [   1  260 1833 ...    0    0    0]]

Maximum review length: 130

```

```

rnn = Sequential()

rnn.add(Embedding(total_words,32,input_length =max_length))
rnn.add(SimpleRNN(64,input_shape = (total_words, max_length), return_sequences=False,activation="relu"))
rnn.add(Dense(1, activation = 'sigmoid')) #flatten

print(rnn.summary())
rnn.compile(loss="binary_crossentropy",optimizer='adam',metrics=["accuracy"])

```

```

→ Model: "sequential"

```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 130, 32)	2964352
simple_rnn (SimpleRNN)	(None, 64)	6208
dense (Dense)	(None, 1)	65

```

=====
Total params: 2,970,625
Trainable params: 2,970,625
Non-trainable params: 0
=====
None

```

```

history = rnn.fit(x_train,y_train,epochs = 20,batch_size=128,verbose = 1)

```

```

→ Epoch 1/20
313/313 [=====] - 96s 286ms/step - loss: 0.6915 - accuracy: 0.5184
Epoch 2/20
313/313 [=====] - 68s 218ms/step - loss: 0.6616 - accuracy: 0.5879
Epoch 3/20
313/313 [=====] - 63s 202ms/step - loss: 0.6626 - accuracy: 0.5762
Epoch 4/20
313/313 [=====] - 56s 180ms/step - loss: 0.5900 - accuracy: 0.6328
Epoch 5/20
313/313 [=====] - 51s 162ms/step - loss: 0.4166 - accuracy: 0.8135

```

```

Epoch 6/20
313/313 [=====] - 51s 161ms/step - loss: 0.3511 - accuracy: 0.8806
Epoch 7/20
313/313 [=====] - 47s 149ms/step - loss: 0.2362 - accuracy: 0.9194
Epoch 8/20
313/313 [=====] - 46s 148ms/step - loss: 0.1676 - accuracy: 0.9421
Epoch 9/20
313/313 [=====] - 45s 144ms/step - loss: 0.3168 - accuracy: 0.8689
Epoch 10/20
313/313 [=====] - 46s 148ms/step - loss: 0.5571 - accuracy: 0.6458
Epoch 11/20
313/313 [=====] - 44s 141ms/step - loss: 0.4791 - accuracy: 0.7574
Epoch 12/20
313/313 [=====] - 45s 144ms/step - loss: 0.2817 - accuracy: 0.9088
Epoch 13/20
313/313 [=====] - 43s 139ms/step - loss: 0.4030 - accuracy: 0.8431
Epoch 14/20
313/313 [=====] - 45s 142ms/step - loss: 0.2630 - accuracy: 0.9154
Epoch 15/20
313/313 [=====] - 43s 138ms/step - loss: 0.2351 - accuracy: 0.9260
Epoch 16/20
313/313 [=====] - 43s 138ms/step - loss: 0.2000 - accuracy: 0.9385
Epoch 17/20
313/313 [=====] - 45s 143ms/step - loss: 0.1599 - accuracy: 0.9498
Epoch 18/20
313/313 [=====] - 44s 140ms/step - loss: 0.1374 - accuracy: 0.9577
Epoch 19/20
313/313 [=====] - 44s 141ms/step - loss: 0.1331 - accuracy: 0.9611
Epoch 20/20
313/313 [=====] - 43s 136ms/step - loss: 0.2814 - accuracy: 0.8869

```

```

model = rnn.save('rnn.h5')
loaded_model = load_model('rnn.h5')

```

```

y_pred = rnn.predict(x_test, batch_size = 128)
print(y_pred)
print(y_test)
for i in range(len(y_pred)):
    if y_pred[i]>0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0

```

```

true = 0
for i, y in enumerate(y_test):
    if y == y_pred[i]:
        true += 1

```

```

print('Correct Prediction: {}'.format(true))
print('Wrong Prediction: {}'.format(len(y_pred) - true))
print('Accuracy: {}'.format(true/len(y_pred)*100))

```

```

➡ 79/79 [=====] - 1s 12ms/step
[[0.78446704]
 [0.02569966]
 [0.78301245]
 ...
 [0.2700789 ]
 [0.72713566]
 [0.78446704]]
2006      1
33575     1
6808      1
32330     0
3777      0
..
40255     1
5864      1
44604     0
42481     1
31671     1
Name: sentiment, Length: 10000, dtype: int64
Correct Prediction: 6918
Wrong Prediction: 3082
Accuracy: 69.17999999999999

```

Message: **Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, everything. I highly recommend it for mystery lovers, for anyone who wants to watch a good movie!**

```
review = str(input('Movie Review: '))
```

```
➡ Movie Review: Nothing was typical about this. Everything was beautifully done in this movie, the story, the flow, the scenario, ever
```

✓ Pre-processing of entered review

```
# Pre-process input
regex = re.compile(r'^a-zA-Z\s')
review = regex.sub('', review)
print('Cleaned: ', review)
```

```
words = review.split(' ')
filtered = [w for w in words if w not in english_stops]
filtered = ' '.join(filtered)
filtered = [filtered.lower()]

print('Filtered: ', filtered)
```

🔄 Cleaned: Nothing was typical about this Everything was beautifully done in this movie the story the flow the scenario everything I
Filtered: ['nothing typical everything beautifully done movie story flow scenario everything i highly recommend mystery lovers any

```
tokenize_words = token.texts_to_sequences(filtered)
tokenize_words = pad_sequences(tokenize_words, maxlen=max_length, padding='post', truncating='post')
print(tokenize_words)
```

🔄

[[76	705	174	1210	126	3	13	2692	2596	174	1	442	280	701
	1771	155	400	33	9	3	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0

✓ Prediction

```
result = rnn.predict(tokenize_words)
print(result)
```

🔄 1/1 [=====] - 0s 40ms/step
[[0.78446704]]

```
if result >= 0.7:
    print('positive')
else:
    print('negative')
```

🔄 positive


Start coding or [generate](#) with AI.

Experiment - 6


✓ Sentiment analysis using RNN-LSTM on tweets data

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/drive/MyDrive/AMITY/Deep Learning (codes)/Data/data.csv')
df.head()
```



	Unnamed: 0	count	hate_speech	offensive_language	neither	class	tweet
0	0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	3	3	0	2	1	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	4	6	0	6	0	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...



```
classes = ['Hate Speech', 'Offensive Language', 'None']
```

```
df.drop(['count', 'hate_speech', 'offensive_language', 'neither', 'Unnamed: 0'], axis=1, inplace=True)
```

```
df.head()
```



	class	tweet
0	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...




```
df.shape
```



```
(24783, 2)
```

```
labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes, counts)
for i in values:
    print(classes[i[0]], ' : ', i[1])
plt.show()
```

```
Hate Speech : 1430
Offensive Language : 19190
None : 4163
```



```
hate_tweets = df[df['class']==0]
offensive_tweets = df[df['class']==1]
neither = df[df['class']==2]
print(hate_tweets.shape)
print(offensive_tweets.shape)
print(neither.shape)
```

```
(1430, 2)
(19190, 2)
(4163, 2)
```

```
for i in range(3):
    hate_tweets = pd.concat([hate_tweets,hate_tweets],ignore_index = True)
neither = pd.concat([neither,neither,neither], ignore_index = True)
offensive_tweets = offensive_tweets.iloc[0:12000,:]
print(hate_tweets.shape)
print(offensive_tweets.shape)
print(neither.shape)
```

```
(11440, 2)
(12000, 2)
(12489, 2)
```

```
df = pd.concat([hate_tweets,offensive_tweets,neither],ignore_index = True)
df.shape
```

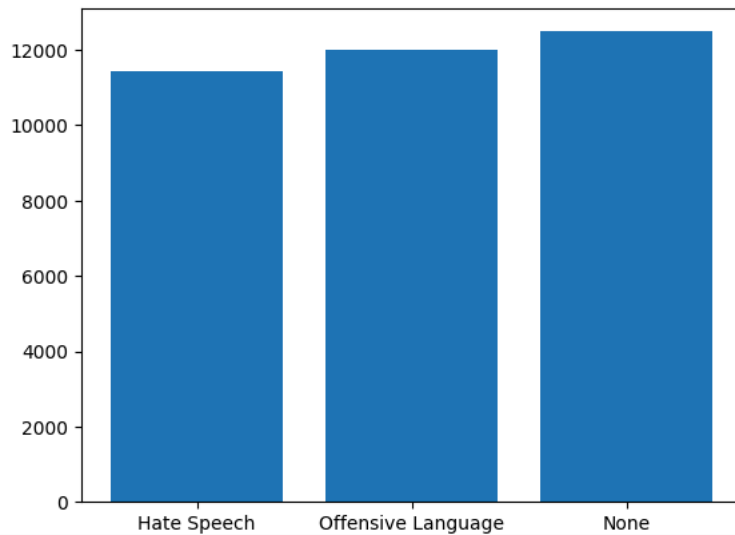
```
(35929, 2)
```

```
labels = df['class']
unique, counts = np.unique(labels, return_counts=True)
values = list(zip(unique, counts))
plt.bar(classes,counts)
for i in values:
    print(classes[i[0]],' : ',i[1])
plt.show()
```

```

Hate Speech : 11440
Offensive Language : 12000
None : 12489

```



```
df.head()
```

```

class tweet
0 0 "@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1..."
1 0 "@CB_Baby24: @white_thunduh alsarabsss" hes a ..."
2 0 "@DevilGrimz: @VigxRArts you're fucking gay, b..."
3 0 "@MarkRoundtreeJr: LMFAOOOO I HATE BLACK PEOPL..."
4 0 "@NoChillPaz: "At least I'm not a nigger" http..."

```

```

import nltk
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import re
nltk.download('wordnet')
nltk.download('stopwords')

```

```

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True

```

```

# dealing with Slangs
d = {'luv':'love', 'wud':'would', 'lyk':'like', 'wateva':'whatever', 'ttyl':'talk to you later',
      'kul':'cool', 'fyn':'fine', 'omg':'oh my god!', 'fam':'family', 'bruh':'brother',
      'cud':'could', 'fud':'food', 'u':'you',
      'ur':'your', 'bday': 'birthday', 'bihday': 'birthday'}

```

```

stop_words = set(stopwords.words("english"))
stop_words.add('rt')
stop_words.remove('not')
lemmatizer = WordNetLemmatizer()
giant_url_regex = ('http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|' '!*\(\),]|(?:%[0-9a-fA-F][0-9a-fA-F]))+')
mention_regex = '@[\w-]+'

```

```

def clean_text(text):
    text = re.sub('""', "", text)
    text = re.sub(mention_regex, ' ', text) #removing all user names
    text = re.sub(giant_url_regex, ' ', text) #removing the urls
    text = text.lower()
    text = re.sub("hm+", "", text) #removing variants of hmmm
    text = re.sub("[^a-z]+", " ", text) #removing all numbers, special chars like @, #, ? etc
    text = text.split()
    text = [word for word in text if not word in stop_words]
    text = [d[word] if word in d else word for word in text] #replacing some slangs
    text = [lemmatizer.lemmatize(token) for token in text]
    text = [lemmatizer.lemmatize(token, "v") for token in text]
    text = " ".join(text)
    return text

```



```
df['processed_tweets'] = df.tweet.apply(lambda x: clean_text(x)) # df.review.map(clean_text) Also can be used
df.head()
```

	class	tweet	processed_tweets
0	0	"@Blackman38Tide: @WhaleLookyHere @HowdyDowdy1...	queer gaywad
1	0	"@CB_Baby24: @white_thunduh alsarabsss" hes a ...	alsarabsss he beaner smh tell he mexican
2	0	"@DevilGrimz: @VigxRArts you're fucking gay, b...	fuck gay blacklist hoe hold tehgodclan anyway
3	0	"@MarkRoundtreeJr: LMFAOOOO I HATE BLACK PEOPL...	lmfaoooo hate black people black people nigger
4	0	"@NoChillPaz: "At least I'm not a nigger" http...	least not nigger lmfaoo

```
x = df.processed_tweets
y = df['class']
print(x.shape)
print(y.shape)
```

```
(35929,)
(35929,)
```

```
# finding unique words
word_unique = []
for i in x:
    for j in i.split():
        word_unique.append(j)
unique, counts = np.unique(word_unique, return_counts=True)
print("The total words in the tweets are : ", len(word_unique))
print("The total UNIQUE words in the tweets are : ", len(unique))
```

```
The total words in the tweets are : 275540
The total UNIQUE words in the tweets are : 14146
```

```
# finding length of tweets
tweets_length = []
for i in x:
    tweets_length.append(len(i.split()))
print("The Average Length tweets are : ", np.mean(tweets_length))
print("The max length of tweets is : ", np.max(tweets_length))
print("The min length of tweets is : ", np.min(tweets_length))
```

```
The Average Length tweets are : 7.669013888502324
The max length of tweets is : 28
The min length of tweets is : 0
```

```
tweets_length = pd.DataFrame(tweets_length)
# tweets_length.describe()
```

	0
count	35929.000000
mean	7.669014
std	3.989625
min	0.000000
25%	4.000000
50%	7.000000
75%	11.000000
max	28.000000

```
#Sorting the Unique words based on their Frequency
col = list(zip(unique, counts))
col = sorted(col, key = lambda x: x[1], reverse=True)
col=pd.DataFrame(col)
print("Top 20 Occuring Words with their frequency are:")
col.iloc[:20,:]
```

→ Top 20 Occuring Words with their frequency are:

	0	1
0	bitch	9066
1	like	3817
2	get	3636
3	hoe	3426
4	trash	3217
5	fuck	3103
6	nigga	2819
7	faggot	2239
8	as	2073
9	you	1851
10	pussy	1847
11	go	1773
12	not	1764
13	bird	1515
14	lol	1494
15	nigger	1459
16	say	1456
17	make	1373
18	amp	1329
19	white	1328

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vectorizer = TfidfVectorizer(max_features = 8000 )  
# tokenize and build vocab
```

```
vectorizer.fit(x)  
# summarize
```

```
print(len(vectorizer.vocabulary_))  
print(vectorizer.idf_.shape)
```

→ 8000
(8000,)

```
x_tfidf = vectorizer.transform(x).toarray()  
print(x_tfidf.shape)
```

→ (35929, 8000)

```
from keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
num_words = 8000  
embed_dim = 32  
tokenizer = Tokenizer(num_words=num_words,oov_token = "<oov>" )  
tokenizer.fit_on_texts(x)  
word_index=tokenizer.word_index  
sequences = tokenizer.texts_to_sequences(x)  
length=[]  
for i in sequences:  
    length.append(len(i))  
print(len(length))  
print("Mean is: ",np.mean(length))  
print("Max is: ",np.max(length))  
print("Min is: ",np.min(length))
```

→ 35929
Mean is: 7.669013888502324
Max is: 28
Min is: 0

```
pad_length = 24
sequences = pad_sequences(sequences, maxlen = pad_length, truncating = 'pre', padding = 'post')
sequences.shape
```

```
(35929, 24)
```

Splitting the Data

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(sequences,y,test_size = 0.05)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(34132, 24)
(1797, 24)
(34132,)
(1797,)
```

RNN Model

```
from keras.layers import Dense, Embedding, Dropout , Activation, Flatten, SimpleRNN
from keras.layers import GlobalMaxPool1D
from keras.models import Model, Sequential
import tensorflow as tf
```

```
recall = tf.keras.metrics.Recall()
precision = tf.keras.metrics.Precision()
```

```
model = Sequential([Embedding(num_words, embed_dim, input_length = pad_length),
                    SimpleRNN(8, return_sequences = True),
                    GlobalMaxPool1D(),
                    Dense(20,activation = 'relu',kernel_initializer='he_uniform'),
                    Dropout(0.25),
                    Dense(3,activation = 'softmax')])
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
# model.name = 'Twitter Hate Text Classification'
model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 24, 32)	256000
simple_rnn_1 (SimpleRNN)	(None, 24, 8)	328
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 8)	0
dense_2 (Dense)	(None, 20)	180
dropout_1 (Dropout)	(None, 20)	0
dense_3 (Dense)	(None, 3)	63
Total params: 256,571		
Trainable params: 256,571		
Non-trainable params: 0		

```
history = model.fit(x = x_train, y = y_train, epochs = 5,validation_split = 0.05)
```

```
Epoch 1/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.2029 - accuracy: 0.9373 - val_loss: 0.1583 - val_accuracy: 0.94
Epoch 2/5
1014/1014 [=====] - 11s 11ms/step - loss: 0.1271 - accuracy: 0.9624 - val_loss: 0.1521 - val_accuracy: 0.94
Epoch 3/5
1014/1014 [=====] - 12s 11ms/step - loss: 0.0906 - accuracy: 0.9754 - val_loss: 0.1479 - val_accuracy: 0.95
Epoch 4/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0714 - accuracy: 0.9801 - val_loss: 0.1644 - val_accuracy: 0.95
Epoch 5/5
1014/1014 [=====] - 12s 12ms/step - loss: 0.0588 - accuracy: 0.9839 - val_loss: 0.1541 - val_accuracy: 0.95
```

```
evaluate = model.evaluate(x_test,y_test)
```

57/57 [=====] - 0s 3ms/step - loss: 0.1112 - accuracy: 0.9688

```
print("Test Acuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))
```

Test Acuracy is : 96.88 %
Test Loss is : 0.1112

```
predictions = model.predict(x_test)
```

57/57 [=====] - 0s 2ms/step

```
predict = []
for i in predictions:
    predict.append(np.argmax(i))
```

```
from sklearn import metrics
cm = metrics.confusion_matrix(predict,y_test)
acc = metrics.accuracy_score(predict,y_test)
```

```
print("The Confusion matrix is: \n",cm)
```

The Confusion matrix is:
[[548 22 1]
[6 572 8]
[0 19 621]]

```
print(acc*100)
```

96.88369504730106

```
from sklearn import metrics
print(metrics.classification_report(y_test, predict))
```

precision recall f1-score support

0	0.96	0.99	0.97	554
1	0.98	0.93	0.95	613
2	0.97	0.99	0.98	630
accuracy			0.97	1797
macro avg	0.97	0.97	0.97	1797
weighted avg	0.97	0.97	0.97	1797

```
from tensorflow.keras.layers import Embedding, LSTM, Dense
# ARCHITECTURE
EMBED_DIM = 32
LSTM_OUT = 64
```

```
model = Sequential()
model.add(Embedding(num_words, EMBED_DIM, input_length = pad_length))
model.add(LSTM(LSTM_OUT))
model.add(Dense(3, activation='softmax'))
model.compile(optimizer = 'adam', loss = 'sparse_categorical_crossentropy', metrics = ['accuracy'])
```

```
print(model.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
embedding_4 (Embedding)	(None, 24, 32)	256000
lstm_1 (LSTM)	(None, 64)	24832
dense_5 (Dense)	(None, 3)	195
=====		
Total params: 281,027		
Trainable params: 281,027		
Non-trainable params: 0		
=====		
None		

```
history = model.fit(x = x_train, y = y_train, epochs = 10,validation_split = 0.05)
```

Epoch 1/10
1014/1014 [=====] - 26s 24ms/step - loss: 0.4328 - accuracy: 0.8246 - val_loss: 0.2531 - val_accuracy: 0.91

```

Epoch 2/10
1014/1014 [=====] - 22s 22ms/step - loss: 0.1874 - accuracy: 0.9405 - val_loss: 0.1876 - val_accuracy: 0.93
Epoch 3/10
1014/1014 [=====] - 23s 23ms/step - loss: 0.1426 - accuracy: 0.9554 - val_loss: 0.1991 - val_accuracy: 0.93
Epoch 4/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.1188 - accuracy: 0.9626 - val_loss: 0.1991 - val_accuracy: 0.93
Epoch 5/10
1014/1014 [=====] - 23s 23ms/step - loss: 0.0999 - accuracy: 0.9688 - val_loss: 0.2142 - val_accuracy: 0.93
Epoch 6/10
1014/1014 [=====] - 23s 22ms/step - loss: 0.0835 - accuracy: 0.9730 - val_loss: 0.1820 - val_accuracy: 0.93
Epoch 7/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0735 - accuracy: 0.9767 - val_loss: 0.2202 - val_accuracy: 0.93
Epoch 8/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0670 - accuracy: 0.9783 - val_loss: 0.2666 - val_accuracy: 0.93
Epoch 9/10
1014/1014 [=====] - 22s 22ms/step - loss: 0.0577 - accuracy: 0.9815 - val_loss: 0.1988 - val_accuracy: 0.93
Epoch 10/10
1014/1014 [=====] - 24s 23ms/step - loss: 0.0537 - accuracy: 0.9823 - val_loss: 0.2313 - val_accuracy: 0.93

```

```
evaluate = model.evaluate(x_test,y_test)
```

```
57/57 [=====] - 0s 8ms/step - loss: 0.1719 - accuracy: 0.9594
```

```

print("Test Acuracy is : {:.2f} %".format(evaluate[1]*100))
print("Test Loss is : {:.4f}".format(evaluate[0]))

```

```

Test Acuracy is : 95.94 %
Test Loss is : 0.1719

```

```
predictions = model.predict(x_test)
```

```
57/57 [=====] - 1s 8ms/step
```

```

predict = []
for i in predictions:
    predict.append(np.argmax(i))

```

✓ Experiment - 7

✓ Auto encoders using MNIST data

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Model

def preprocess(array):
    """
    Normalizes the supplied array and reshapes it into the appropriate format.
    """

    array = array.astype("float32") / 255.0
    array = np.reshape(array, (len(array), 28, 28, 1))
    return array

def noise(array):
    """
    Adds random noise to each image in the supplied array.
    """

    noise_factor = 0.4 #amount of noise to add
    noisy_array = array + noise_factor * np.random.normal(
        loc=0.0, scale=1.0, size=array.shape
    )

    return np.clip(noisy_array, 0.0, 1.0)

def display(array1, array2):
    """
    Displays ten random images from each one of the supplied arrays.
    """

    n = 10

    indices = np.random.randint(len(array1), size=n)
    images1 = array1[indices, :]
    images2 = array2[indices, :]

    plt.figure(figsize=(20, 4))
    for i, (image1, image2) in enumerate(zip(images1, images2)):
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(image1.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(image2.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

    plt.show()

# Since we only need images from the dataset to encode and decode, we
# won't use the labels.
(train_data, _), (test_data, _) = mnist.load_data()


# Normalize and reshape the data
train_data = preprocess(train_data)
test_data = preprocess(test_data)

# Create a copy of the data with added noise
noisy_train_data = noise(train_data)
noisy_test_data = noise(test_data)

# Display the train data and a version of it with added noise
display(train_data, noisy_train_data)
```

display(train_data, noisy_train_data,

→ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 ————— 2s 0us/step



```
input = layers.Input(shape=(28, 28, 1))
```

```
# Encoder
```

```
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(input)
```

```
x = layers.MaxPooling2D((2, 2), padding="same")(x)
```

```
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(x)
```

```
x = layers.MaxPooling2D((2, 2), padding="same")(x)
```

```
# Decoder
```

```
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
```

```
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
```

```
x = layers.Conv2D(1, (3, 3), activation="sigmoid", padding="same")(x)
```

```
# Autoencoder
```

```
autoencoder = Model(input, x)
```

```
autoencoder.compile(optimizer="adam", loss="binary_crossentropy")
```

```
autoencoder.summary()
```

→ Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 28, 28, 1)	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_transpose (Conv2DTranspose)	(None, 14, 14, 32)	9,248
conv2d_transpose_1 (Conv2DTranspose)	(None, 28, 28, 32)	9,248
conv2d_2 (Conv2D)	(None, 28, 28, 1)	289

Total params: 28,353 (110.75 KB)

```
autoencoder.fit(  
    x=train_data,  
    y=train_data,  
    epochs=50,  
    batch_size=128,  
    shuffle=True,  
    validation_data=(test_data, test_data),  
)
```

→ Epoch 1/50
469/469 ————— 10s 12ms/step - loss: 0.2439 - val_loss: 0.0743

Epoch 2/50
469/469 ————— 3s 6ms/step - loss: 0.0736 - val_loss: 0.0701

Epoch 3/50
469/469 ————— 3s 6ms/step - loss: 0.0704 - val_loss: 0.0685

Epoch 4/50
469/469 ————— 3s 6ms/step - loss: 0.0689 - val_loss: 0.0676

Epoch 5/50
469/469 ————— 3s 6ms/step - loss: 0.0679 - val_loss: 0.0669

Epoch 6/50
469/469 ————— 3s 6ms/step - loss: 0.0672 - val_loss: 0.0663


Epoch 7/50
469/469 ————— 5s 6ms/step - loss: 0.0667 - val_loss: 0.0659

Epoch 8/50
469/469 ————— 5s 6ms/step - loss: 0.0663 - val_loss: 0.0656

Epoch 9/50
469/469 ————— 3s 6ms/step - loss: 0.0659 - val_loss: 0.0652
Epoch 10/50
469/469 ————— 5s 6ms/step - loss: 0.0656 - val_loss: 0.0649
Epoch 11/50
469/469 ————— 5s 6ms/step - loss: 0.0654 - val_loss: 0.0647
Epoch 12/50
469/469 ————— 5s 6ms/step - loss: 0.0650 - val_loss: 0.0644
Epoch 13/50
469/469 ————— 5s 6ms/step - loss: 0.0648 - val_loss: 0.0644
Epoch 14/50
469/469 ————— 5s 6ms/step - loss: 0.0645 - val_loss: 0.0641
Epoch 15/50
469/469 ————— 5s 7ms/step - loss: 0.0643 - val_loss: 0.0639
Epoch 16/50
469/469 ————— 5s 6ms/step - loss: 0.0642 - val_loss: 0.0637
Epoch 17/50
469/469 ————— 3s 6ms/step - loss: 0.0641 - val_loss: 0.0636
Epoch 18/50
469/469 ————— 5s 6ms/step - loss: 0.0640 - val_loss: 0.0634
Epoch 19/50
469/469 ————— 5s 6ms/step - loss: 0.0638 - val_loss: 0.0634
Epoch 20/50
469/469 ————— 3s 6ms/step - loss: 0.0637 - val_loss: 0.0633
Epoch 21/50
469/469 ————— 3s 7ms/step - loss: 0.0636 - val_loss: 0.0632
Epoch 22/50
469/469 ————— 5s 6ms/step - loss: 0.0634 - val_loss: 0.0631
Epoch 23/50
469/469 ————— 5s 6ms/step - loss: 0.0635 - val_loss: 0.0630
Epoch 24/50
469/469 ————— 3s 7ms/step - loss: 0.0634 - val_loss: 0.0629
Epoch 25/50
469/469 ————— 3s 6ms/step - loss: 0.0634 - val_loss: 0.0629
Epoch 26/50
469/469 ————— 3s 6ms/step - loss: 0.0633 - val_loss: 0.0629
Epoch 27/50
469/469 ————— 5s 7ms/step - loss: 0.0632 - val_loss: 0.0628
Epoch 28/50
469/469 ————— 3s 6ms/step - loss: 0.0630 - val_loss: 0.0627
Epoch 29/50
469/469 ————— 3s 6ms/step - loss: 0.0630 - val_loss: 0.0626

```
predictions = autoencoder.predict(test_data)
display(test_data, predictions)
```

↔ 313/313 ————— 1s 2ms/step



Experiment - 8

✓ variational-autoencoders

```
# For working with and visualizing the data
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
# For training the VAE
import tensorflow as tf
```

```
# For creating interactive widgets
import ipywidgets as widgets
from IPython.display import display
```

2024-05-22 06:28:22.848294: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in the environment. To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
# Load the data from a .csv file
pixel_data = pd.read_csv('/workspace/Bootcamp/data/age_gender.csv')['pixels']
```

```
# Shuffle the data
pixel_data = pixel_data.sample(frac=1.0, random_state=1)
# Convert the data into a NumPy array
pixel_data = pixel_data.apply(lambda x: np.array(x.split(" "), dtype=np.int))
pixel_data = np.stack(np.array(pixel_data), axis=0)
# Rescale pixel values to be between 0 and 1
pixel_data = pixel_data * (1./255)
```

/tmp/ipykernel_92174/3099532490.py:4: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` instead of `np.int` in the future. This has no effect for your code. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
pixel_data = pixel_data.apply(lambda x: np.array(x.split(" "), dtype=np.int))
```

```
# The data is now a NumPy array of 23705 images.
# we are working with 48x48x1 images)
pixel_data.shape
```

(23705, 2304)

```
class Sampling(tf.keras.layers.Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs # Unpack the inputs into mean and log-variance
        batch = tf.shape(z_mean)[0] # Get the batch size
        dim = tf.shape(z_mean)[1] # Get the dimensionality of the latent space
        epsilon = tf.keras.backend.random_normal(shape=(batch, dim)) # Sample from standard normal distribution
        return epsilon * tf.exp(z_log_var * 0.5) + z_mean # Apply the reparameterization trick
```

```
def build_vae(num_pixels, num_latent_vars=3):
```

```
    # Encoder
    encoder_inputs = tf.keras.Input(shape=(num_pixels,)) # Input layer for the encoder
    x = tf.keras.layers.Dense(512, activation='relu')(encoder_inputs) # First dense layer with 512 units and ReLU activation
    x = tf.keras.layers.Dense(128, activation='relu')(x) # Second dense layer with 128 units and ReLU activation
    x = tf.keras.layers.Dense(32, activation='relu')(x) # Third dense layer with 32 units and ReLU activation
    z_mean = tf.keras.layers.Dense(num_latent_vars)(x) # Dense layer for the mean of the latent variables
    z_log_var = tf.keras.layers.Dense(num_latent_vars)(z_mean) # Dense layer for the log-variance of the latent variables
    z = Sampling()(z_mean, z_log_var) # Sampling layer to sample the latent variables using the reparameterization trick
```

```
    encoder = tf.keras.Model(inputs=encoder_inputs, outputs=z) # Define the encoder model
```

```
    # Decoder
    decoder_inputs = tf.keras.Input(shape=(num_latent_vars,)) # Input layer for the decoder
    x = tf.keras.layers.Dense(32, activation='relu')(decoder_inputs) # First dense layer with 32 units and ReLU activation
    x = tf.keras.layers.Dense(128, activation='relu')(x) # Second dense layer with 128 units and ReLU activation
    x = tf.keras.layers.Dense(512, activation='relu')(x) # Third dense layer with 512 units and ReLU activation
    reconstruction = tf.keras.layers.Dense(num_pixels, activation='linear')(x) # Output dense layer with 'num_pixels' units and linear activation
```

```
    decoder = tf.keras.Model(inputs=decoder_inputs, outputs=reconstruction) # Define the decoder model
```

```
    # Full model
```

```
model_inputs = encoder.input # Inputs of the full VAE model are the inputs of the encoder
model_outputs = decoder(encoder.output) # Outputs of the full VAE model are the outputs of the decoder, given the encoder's output

model = tf.keras.Model(inputs=model_inputs, outputs=model_outputs) # Define the full VAE model

# Compile model for training
model.compile(
    optimizer='adam', # Adam optimizer
    loss='mse' # Mean Squared Error (MSE) loss function
)

# Return all three models
return encoder, decoder, model # Return the encoder, decoder, and full VAE models
```

```
face_encoder, face_decoder, face_model = build_vae(num_pixels=2304, num_latent_vars=3)
```

```
2024-05-22 06:28:31.067844: W tensorflow/core/common_runtime/gpu/gpu_device.cc:1960] Cannot dlopen some GPU libraries. Please make :
Skipping registering GPU devices...
```

```
print(face_encoder.summary())
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 2304)]	0	[]
dense (Dense)	(None, 512)	1180160	['input_1[0][0]']
dense_1 (Dense)	(None, 128)	65664	['dense[0][0]']
dense_2 (Dense)	(None, 32)	4128	['dense_1[0][0]']
dense_3 (Dense)	(None, 3)	99	['dense_2[0][0]']
dense_4 (Dense)	(None, 3)	12	['dense_3[0][0]']
sampling (Sampling)	(None, 3)	0	['dense_3[0][0]', 'dense_4[0][0]']

Total params: 1250063 (4.77 MB)
Trainable params: 1250063 (4.77 MB)
Non-trainable params: 0 (0.00 Byte)

None

```
print(face_decoder.summary())
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 3)]	0
dense_5 (Dense)	(None, 32)	128
dense_6 (Dense)	(None, 128)	4224
dense_7 (Dense)	(None, 512)	66048
dense_8 (Dense)	(None, 2304)	1181952

Total params: 1252352 (4.78 MB)
Trainable params: 1252352 (4.78 MB)
Non-trainable params: 0 (0.00 Byte)

None

```
history = face_model.fit(
    pixel_data,
    pixel_data,
    validation_split=0.2,
    batch_size=32,
    epochs=100,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(
            monitor='val_loss',
            patience=10,
            restore_best_weights=True
        )
    ]
)
```

```
]
)
```

```
Epoch 1/100
593/593 [=====] - 8s 12ms/step - loss: 0.0256 - val_loss: 0.0216
Epoch 2/100
593/593 [=====] - 5s 9ms/step - loss: 0.0216 - val_loss: 0.0213
Epoch 3/100
593/593 [=====] - 5s 9ms/step - loss: 0.0214 - val_loss: 0.0212
Epoch 4/100
593/593 [=====] - 5s 9ms/step - loss: 0.0212 - val_loss: 0.0212
Epoch 5/100
593/593 [=====] - 5s 9ms/step - loss: 0.0212 - val_loss: 0.0210
Epoch 6/100
593/593 [=====] - 5s 9ms/step - loss: 0.0211 - val_loss: 0.0213
Epoch 7/100
593/593 [=====] - 5s 9ms/step - loss: 0.0210 - val_loss: 0.0209
Epoch 8/100
593/593 [=====] - 5s 9ms/step - loss: 0.0209 - val_loss: 0.0209
Epoch 9/100
593/593 [=====] - 5s 9ms/step - loss: 0.0209 - val_loss: 0.0210
Epoch 10/100
593/593 [=====] - 5s 9ms/step - loss: 0.0208 - val_loss: 0.0206
Epoch 11/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 - val_loss: 0.0206
Epoch 12/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 - val_loss: 0.0208
Epoch 13/100
593/593 [=====] - 6s 9ms/step - loss: 0.0206 - val_loss: 0.0207
Epoch 14/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 - val_loss: 0.0205
Epoch 15/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0210
Epoch 16/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0204
Epoch 17/100
593/593 [=====] - 5s 9ms/step - loss: 0.0208 - val_loss: 0.0208
Epoch 18/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 - val_loss: 0.0206
Epoch 19/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 - val_loss: 0.0207
Epoch 20/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0205
Epoch 21/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0206
Epoch 22/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0209
Epoch 23/100
593/593 [=====] - 5s 9ms/step - loss: 0.0207 - val_loss: 0.0207
Epoch 24/100
593/593 [=====] - 5s 9ms/step - loss: 0.0206 - val_loss: 0.0210
Epoch 25/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 - val_loss: 0.0205
Epoch 26/100
593/593 [=====] - 5s 9ms/step - loss: 0.0205 - val_loss: 0.0205
```

```
i = 6
```

```
sample = np.array(pixel_data)[i].copy()
sample = sample.reshape(48, 48, 1)
```

```
reconstruction = face_model.predict(pixel_data)[i].copy()
reconstruction = reconstruction.reshape(48, 48, 1)
```

```
plt.figure(figsize=(10, 5))
```

```
plt.subplot(1, 2, 1)
plt.imshow(sample, cmap='gray')
plt.axis('off')
plt.title("Original Image")
```

```
plt.subplot(1, 2, 2)
plt.imshow(reconstruction, cmap='gray')
plt.axis('off')
plt.title("Reconstructed Image")
```

```
plt.show()
```

Original Image



Reconstructed Image



A function to allow us to specify our own latent variable values and plot the constructed image

```
def generate_face_image(latent1, latent2, latent3):
    latent_vars = np.array([[latent1, latent2, latent3]])
    reconstruction = np.array(face_decoder(latent_vars))
    reconstruction = reconstruction.reshape(48, 48, 1)
    plt.figure()
    plt.imshow(reconstruction, cmap='gray')
    plt.axis('off')
    plt.show()
```

Let's get the min and max for each slider on the interactive widget

```
latent1_min = np.min(face_encoder(pixel_data).numpy()[0, 0])
latent1_max = np.max(face_encoder(pixel_data).numpy()[0, 0])
```

```
latent2_min = np.min(face_encoder(pixel_data).numpy()[0, 1])
latent2_max = np.max(face_encoder(pixel_data).numpy()[0, 1])
```

```
latent3_min = np.min(face_encoder(pixel_data).numpy()[0, 2])
latent3_max = np.max(face_encoder(pixel_data).numpy()[0, 2])
```

```
import tensorflow as tf
print(tf.__version__)
```


2.13.1

Create the interactive widget

```
face_image_generator = widgets.interact(
    generate_face_image,
    latent1=(latent1_min, latent1_max),
    latent2=(latent2_min, latent2_max),
    latent3=(latent3_min, latent3_max),
)
```

Display the widget

```
display(face_image_generator)
```



Experiment - 9

✓ Implementing GAN Architecture on MINST dataset

```
from tensorflow.keras.layers import (Dense,
                                     BatchNormalization,
                                     LeakyReLU,
                                     Reshape,
                                     Conv2DTranspose,
                                     Conv2D,
                                     Dropout,
                                     Flatten)

import tensorflow as tf
import matplotlib.pyplot as plt

# underscore to omit the label arrays
(train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()

train_images = train_images.reshape(train_images.shape[0], 28, 28, 1).astype('float32')
train_images = (train_images - 127.5) / 127.5 # Normalize the images to [-1, 1]

BUFFER_SIZE = 60000
BATCH_SIZE = 256

# Batch and shuffle the data
train_dataset = tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE).batch(BATCH_SIZE)

➡ Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

# Set the dimensions of the noise
z_dim = 100

def generator_model():
    model = tf.keras.Sequential()
    model.add(Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256) # Note: None is the batch size

    model.add(Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', use_bias=False))
    assert model.output_shape == (None, 14, 14, 64)
    model.add(BatchNormalization())
    model.add(LeakyReLU())

    model.add(Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same', use_bias=False, activation='tanh'))
    assert model.output_shape == (None, 28, 28, 1)

    print(model.summary())

    return model

generator = generator_model()
```

➡ Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 12544)	1254400
batch_normalization (Batch Normalization)	(None, 12544)	50176
leaky_re_lu (LeakyReLU)	(None, 12544)	0
reshape (Reshape)	(None, 7, 7, 256)	0

conv2d_transpose (Conv2DTra nspose)	(None, 7, 7, 128)	819200
batch_normalization_1 (Batc hNormalization)	(None, 7, 7, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 7, 7, 128)	0
conv2d_transpose_1 (Conv2DT ranspose)	(None, 14, 14, 64)	204800
batch_normalization_2 (Batc hNormalization)	(None, 14, 14, 64)	256
leaky_re_lu_2 (LeakyReLU)	(None, 14, 14, 64)	0
conv2d_transpose_2 (Conv2DT ranspose)	(None, 28, 28, 1)	1600

```

=====
Total params: 2,330,944
Trainable params: 2,305,472
Non-trainable params: 25,472


```

None

```

# Create a random noise and generate a sample
noise = tf.random.normal([1, 100])
generated_image = generator(noise, training=False)
# Visualize the generated sample
plt.imshow(generated_image[0, :, :, 0], cmap='gray')

```



```

def discriminator_model():
    model = tf.keras.Sequential()

    model.add(Conv2D(64, (5, 5), strides=(2, 2), padding='same', input_shape=[28, 28, 1]))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))

    model.add(Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(LeakyReLU())
    model.add(Dropout(0.3))

    model.add(Flatten())
    model.add(Dense(1))

    print(model.summary())

    return model

```

discriminator = discriminator_model()

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 14, 14, 64)	1664
leaky_re_lu_4 (LeakyReLU)	(None, 14, 14, 64)	0

dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 128)	204928
leaky_re_lu_5 (LeakyReLU)	(None, 7, 7, 128)	0
dropout_1 (Dropout)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 1)	6273

=====

Total params: 212,865
Trainable params: 212,865
Non-trainable params: 0

None

```
# This method returns a helper function to compute cross entropy loss
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    total_loss = real_loss + fake_loss
    return total_loss

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)

generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)

import os
checkpoint_dir = '/content/drive/MyDrive/AMITY/Deep Learning (codes)/GAN/'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                  discriminator_optimizer=discriminator_optimizer,
                                  generator=generator,
                                  discriminator=discriminator)

EPOCHS = 60
# We will reuse this seed overtime (so it's easier)
# to visualize progress in the animated GIF
num_examples_to_generate = 16
noise_dim = 100
seed = tf.random.normal([num_examples_to_generate, noise_dim])

# tf.function annotation causes the function
# to be "compiled" as part of the training
@tf.function
def train_step(images):

    # 1 - Create a random noise to feed it into the model
    # for the image generation
    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    # 2 - Generate images and calculate loss values
    # GradientTape method records operations for automatic differentiation.
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

    # 3 - Calculate gradients using loss values and model variables
    # "gradient" method computes the gradient using
    # operations recorded in context of this tape (gen_tape and disc_tape).

    # It accepts a target (e.g., gen_loss) variable and
    # a source variable (e.g., generator.trainable_variables)
    # target --> a list or nested structure of Tensors or Variables to be differentiated.
    # source --> a list or nested structure of Tensors or Variables.
    # target will be differentiated against elements in sources.

    # "gradient" method returns a list or nested structure of Tensors
```


[illegible]



```
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))
```



```
<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x7fa5c2d49450>
```

```
# PIL is a library which may open different image file formats
import PIL
# Display a single image using the epoch number
def display_image(epoch_no):
    return PIL.Image.open('image_at_epoch_{:04d}.png'.format(epoch_no))
display_image(EPOCHS)
```



```
import glob # The glob module is used for Unix style pathname pattern expansion.
import imageio # The library that provides an easy interface to read and write a wide range of image data
```

```
anim_file = 'dcgan.gif'
```

```
with imageio.get_writer(anim_file, mode='I') as writer:
    filenames = glob.glob('image*.png')
    filenames = sorted(filenames)
    for filename in filenames:
        image = imageio.imread(filename)
        writer.append_data(image)
    # image = imageio.imread(filename)
    # writer.append_data(image)
```

```
display.Image(open('dcgan.gif', 'rb').read())
```